



BACCHANAL BUFFET

A COMPREHENSIVE STUDY
USING YELP DATA

AHMET FARUK İZGÖRDÜ

Table Of Content

- Introduction
- Methodology
- Results
- Conclusion

● Introduction

Overview of Business Need

This project addresses the need for a systematic approach to assess customer sentiments from Yelp reviews, enabling the restaurant to make data-driven decisions that enhance customer satisfaction and loyalty.

Solution Approach

To address the business need, we have implemented a sentiment analysis model on Yelp reviews using advanced natural language processing (NLP) techniques.

Objectives

Analyze Customer Sentiments: Overall sentiment of customer reviews and identify key areas of customer satisfaction and dissatisfaction.

Visualize Data: Create clear and informative visualizations to present sentiment trends and patterns

Provide Actionable Insights: Offer recommendations based on sentiment analysis



● Introduction

Bacchanal Buffet Claimed

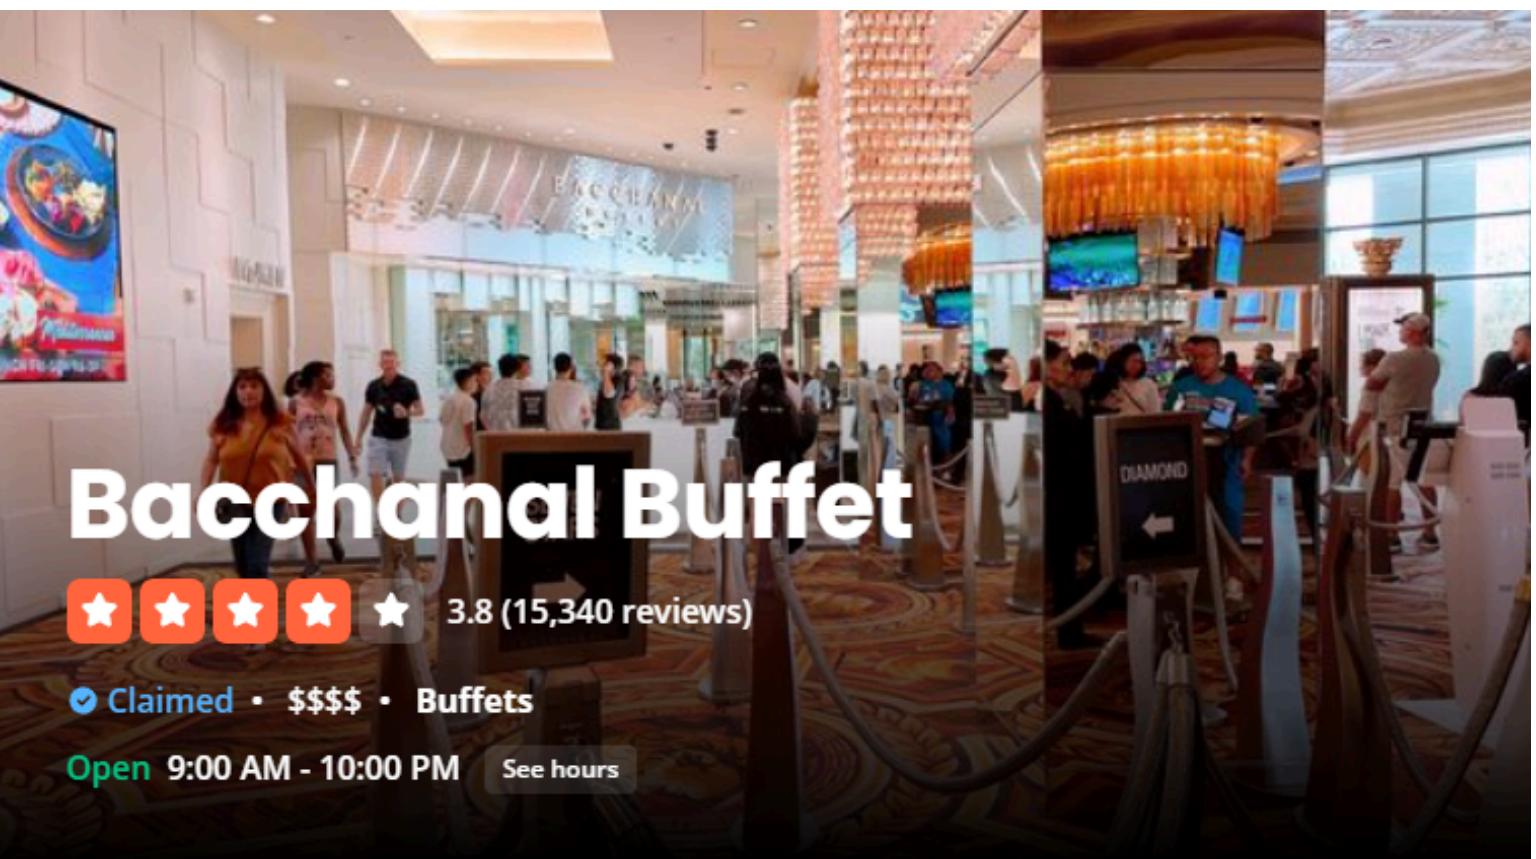
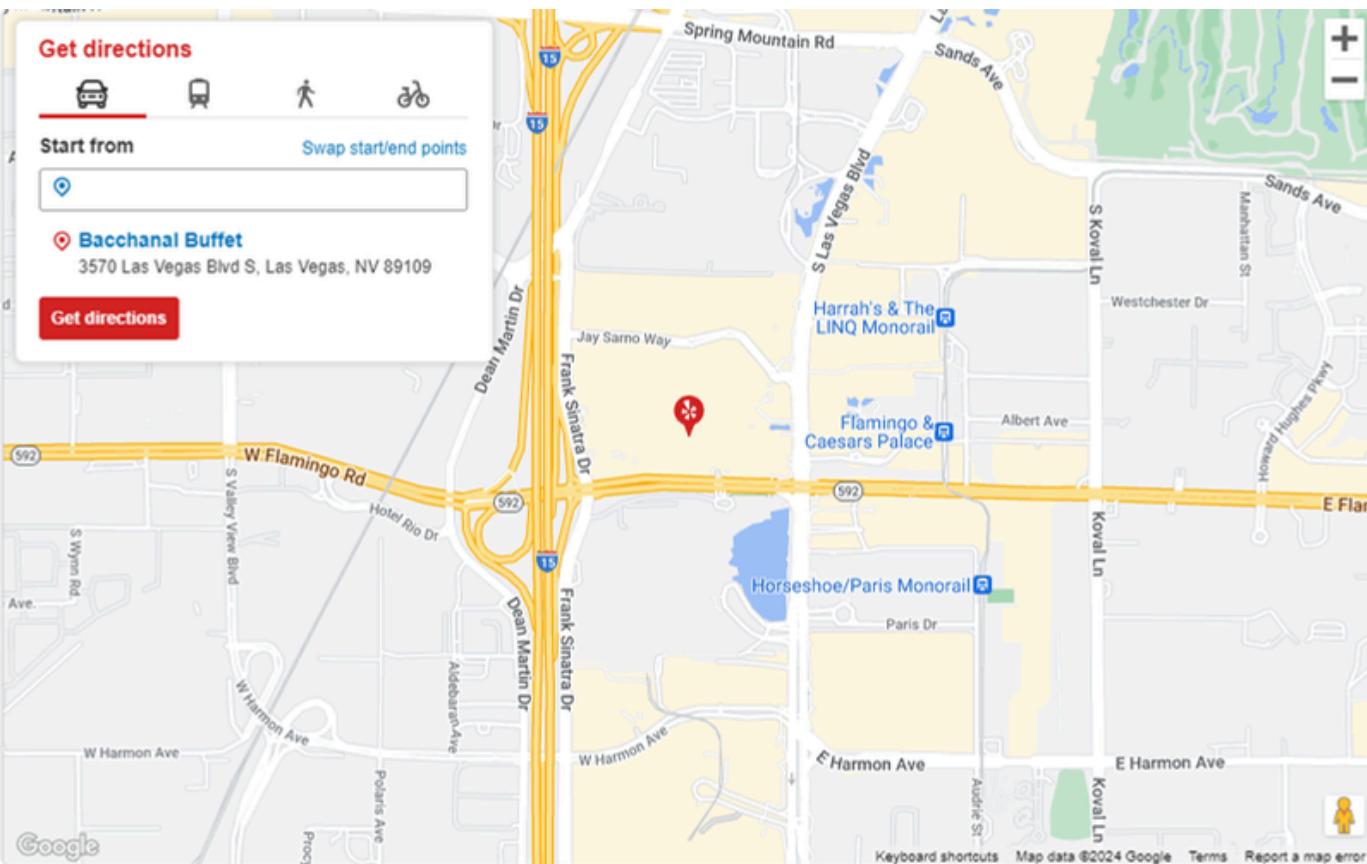
6,493 reviews NEW AI Reviews Summary #106 of 540 Quick Bites in Las Vegas | \$\$ - \$\$\$, Quick Bites, Italian, American

3570 Las Vegas Boulevard South Caesars Palace, Las Vegas, NV 89109 | +1 702-731-7928 | Website | Menu | Open now 9:00 am - 10:00 pm | Improve this listing

Reserve a table OpenTable

Mon, 29 July 7:00 pm 2 21:00 21:15 21:30 21:45

See all (3484)



Bacchanal Buffet

4.8 (15,340 reviews) 3.8 (15,340 reviews)

Claimed • \$\$\$ • Buffets

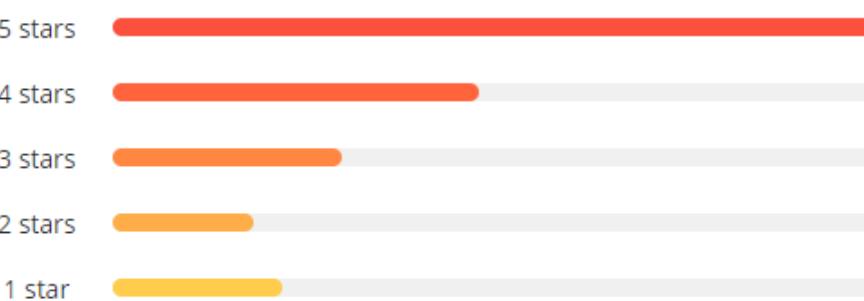
Open 9:00 AM - 10:00 PM See hours

THE LARGEST BUFFET IN LAS VEGAS

Over 25,000 square feet to indulge on flavors from around the globe.

10 KITCHENS 9 CHEF-ATTENDED ACTION STATIONS 250+ MENU ITEMS ∞ INFINITE FLAVORS

Overall rating
4.8 stars
15,340 reviews



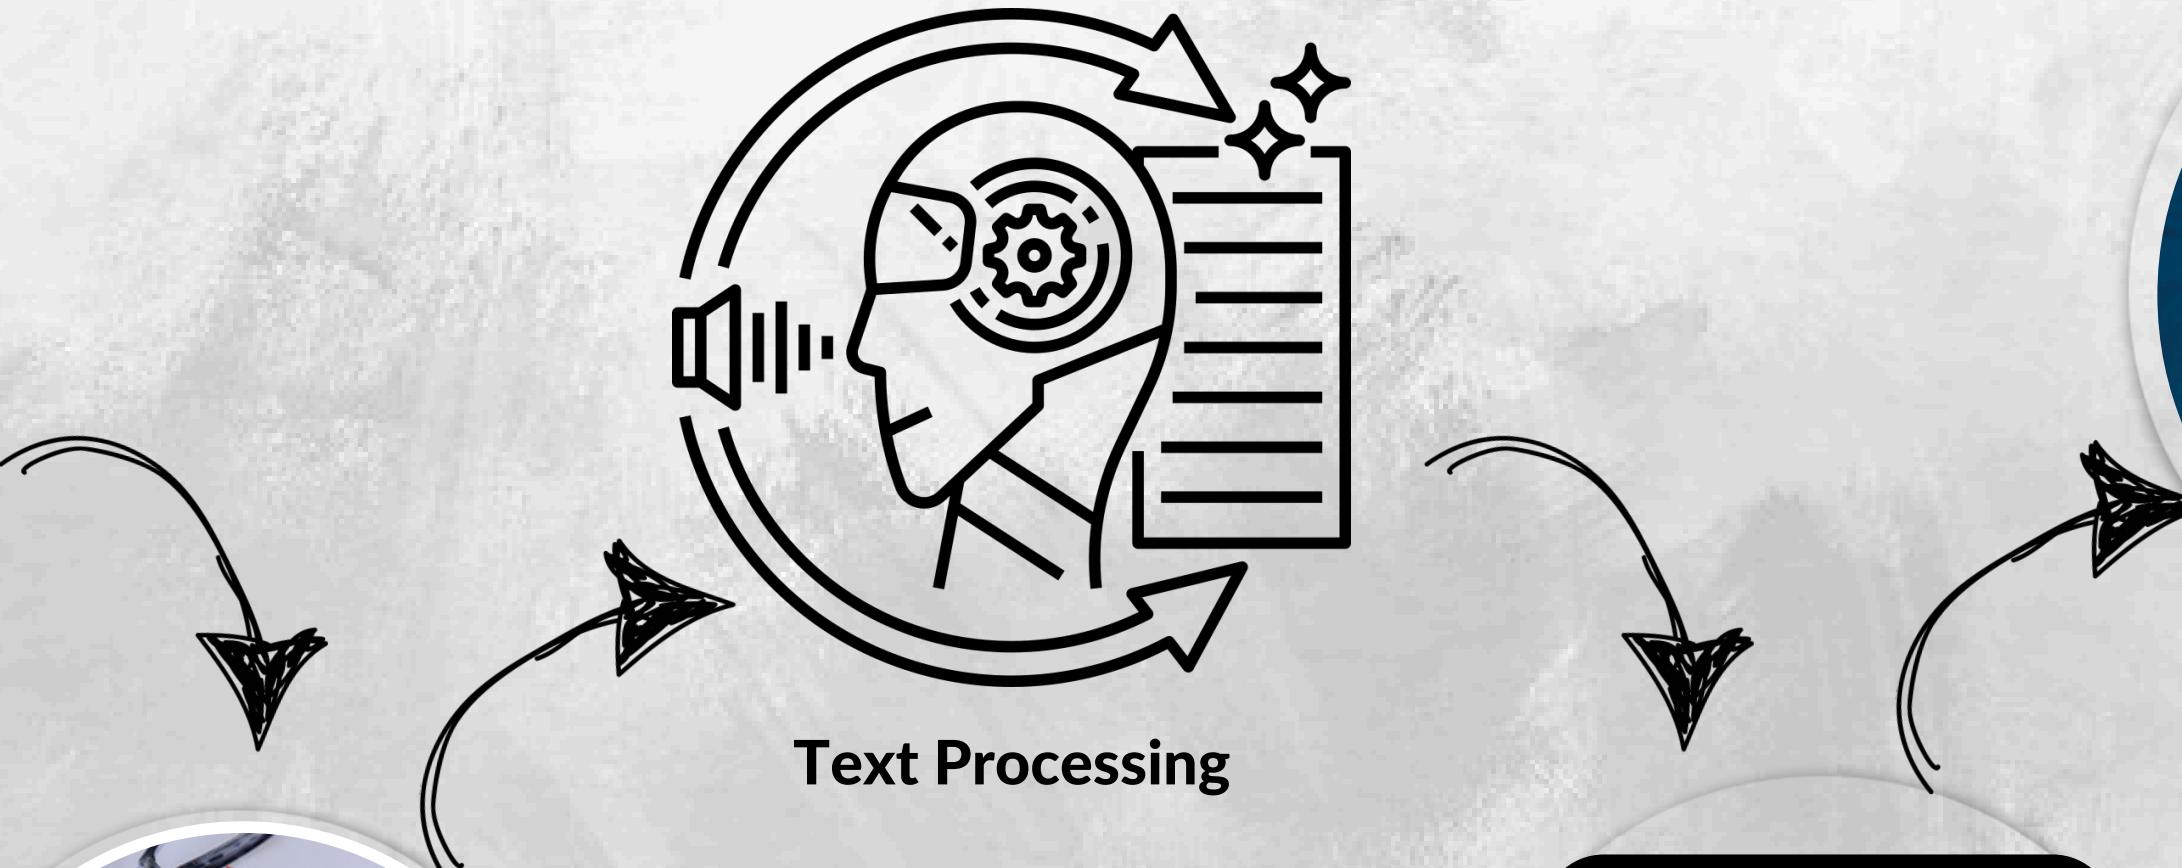
Methodology



Data Acquire



Exploratory
Data Analysis



Modelling



Evaluation

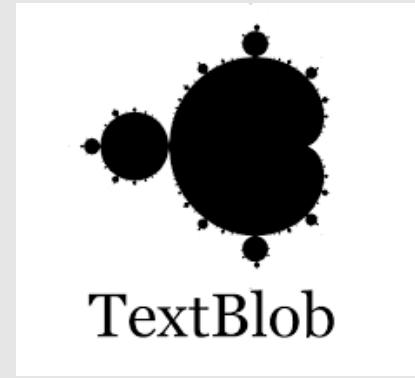
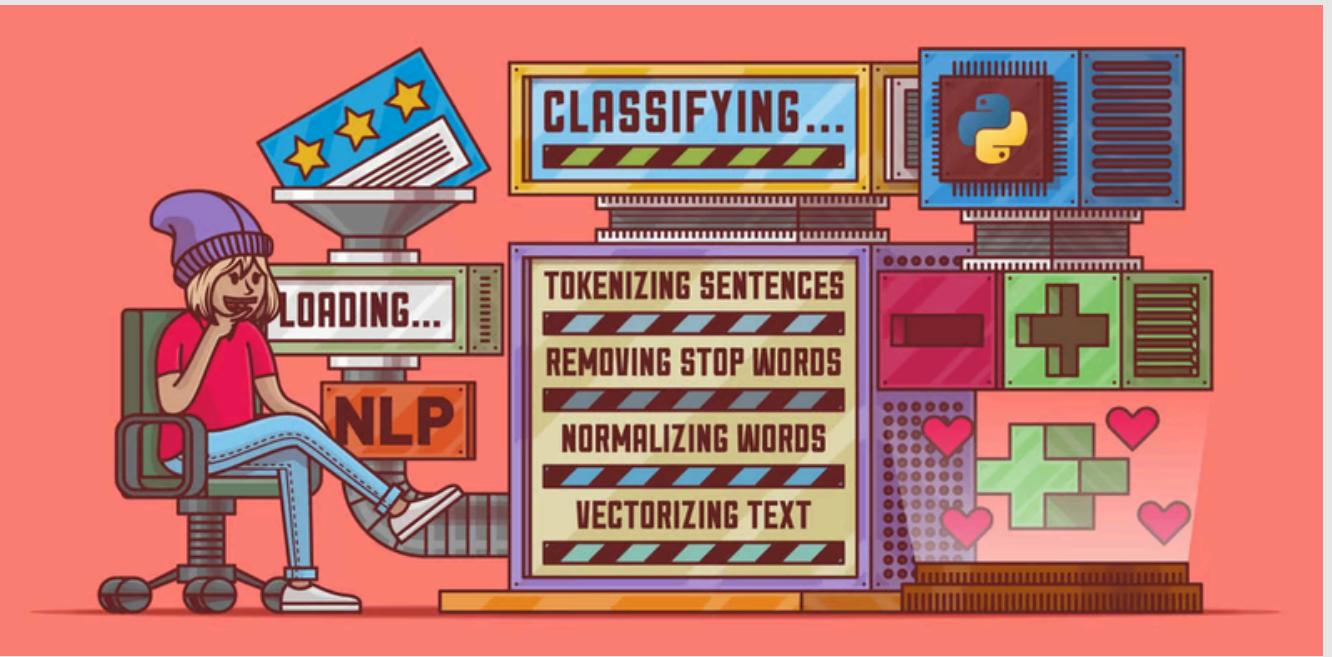
Methodology



Data Acquire



Methodology



TextBlob



NLTK



wordcloud



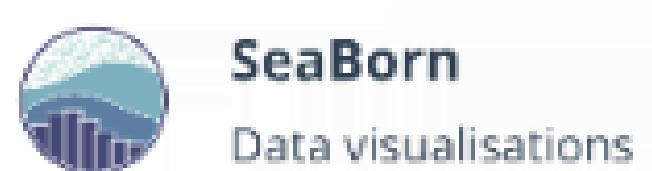
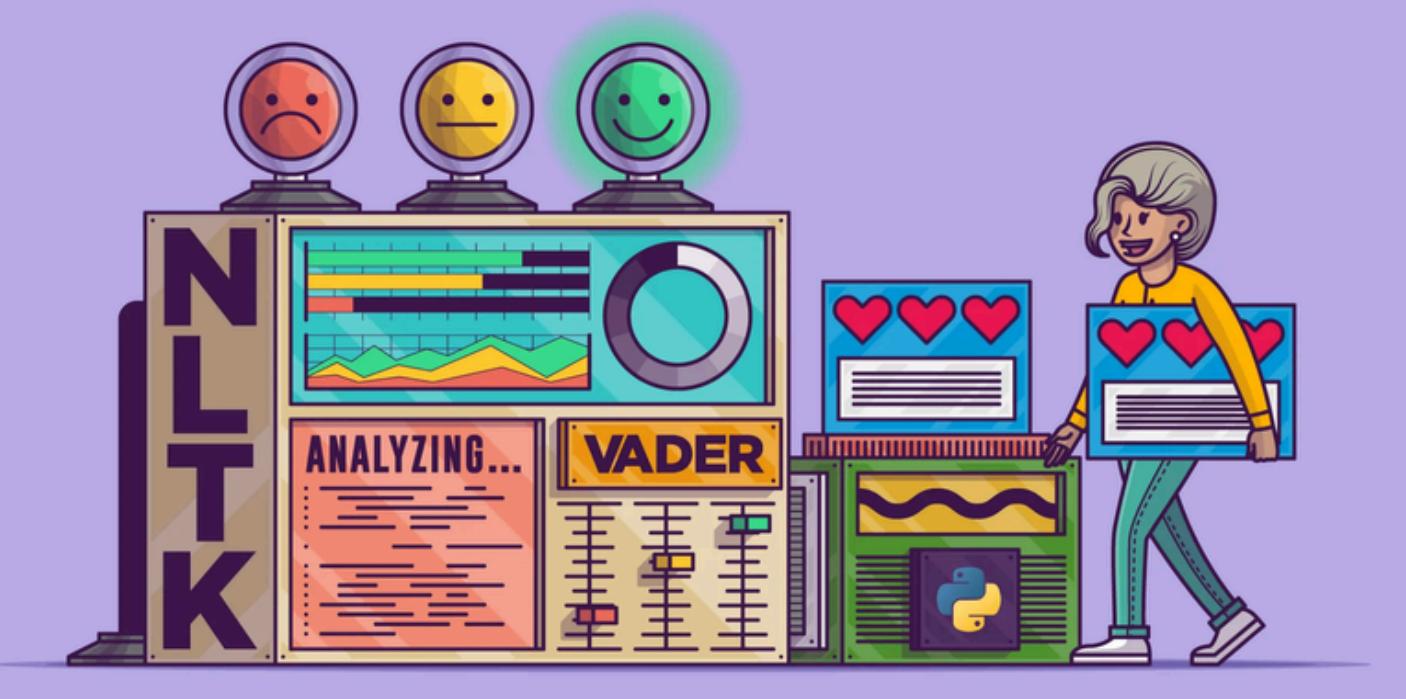
TOOLS



NumPy
Mathematical functions



python™



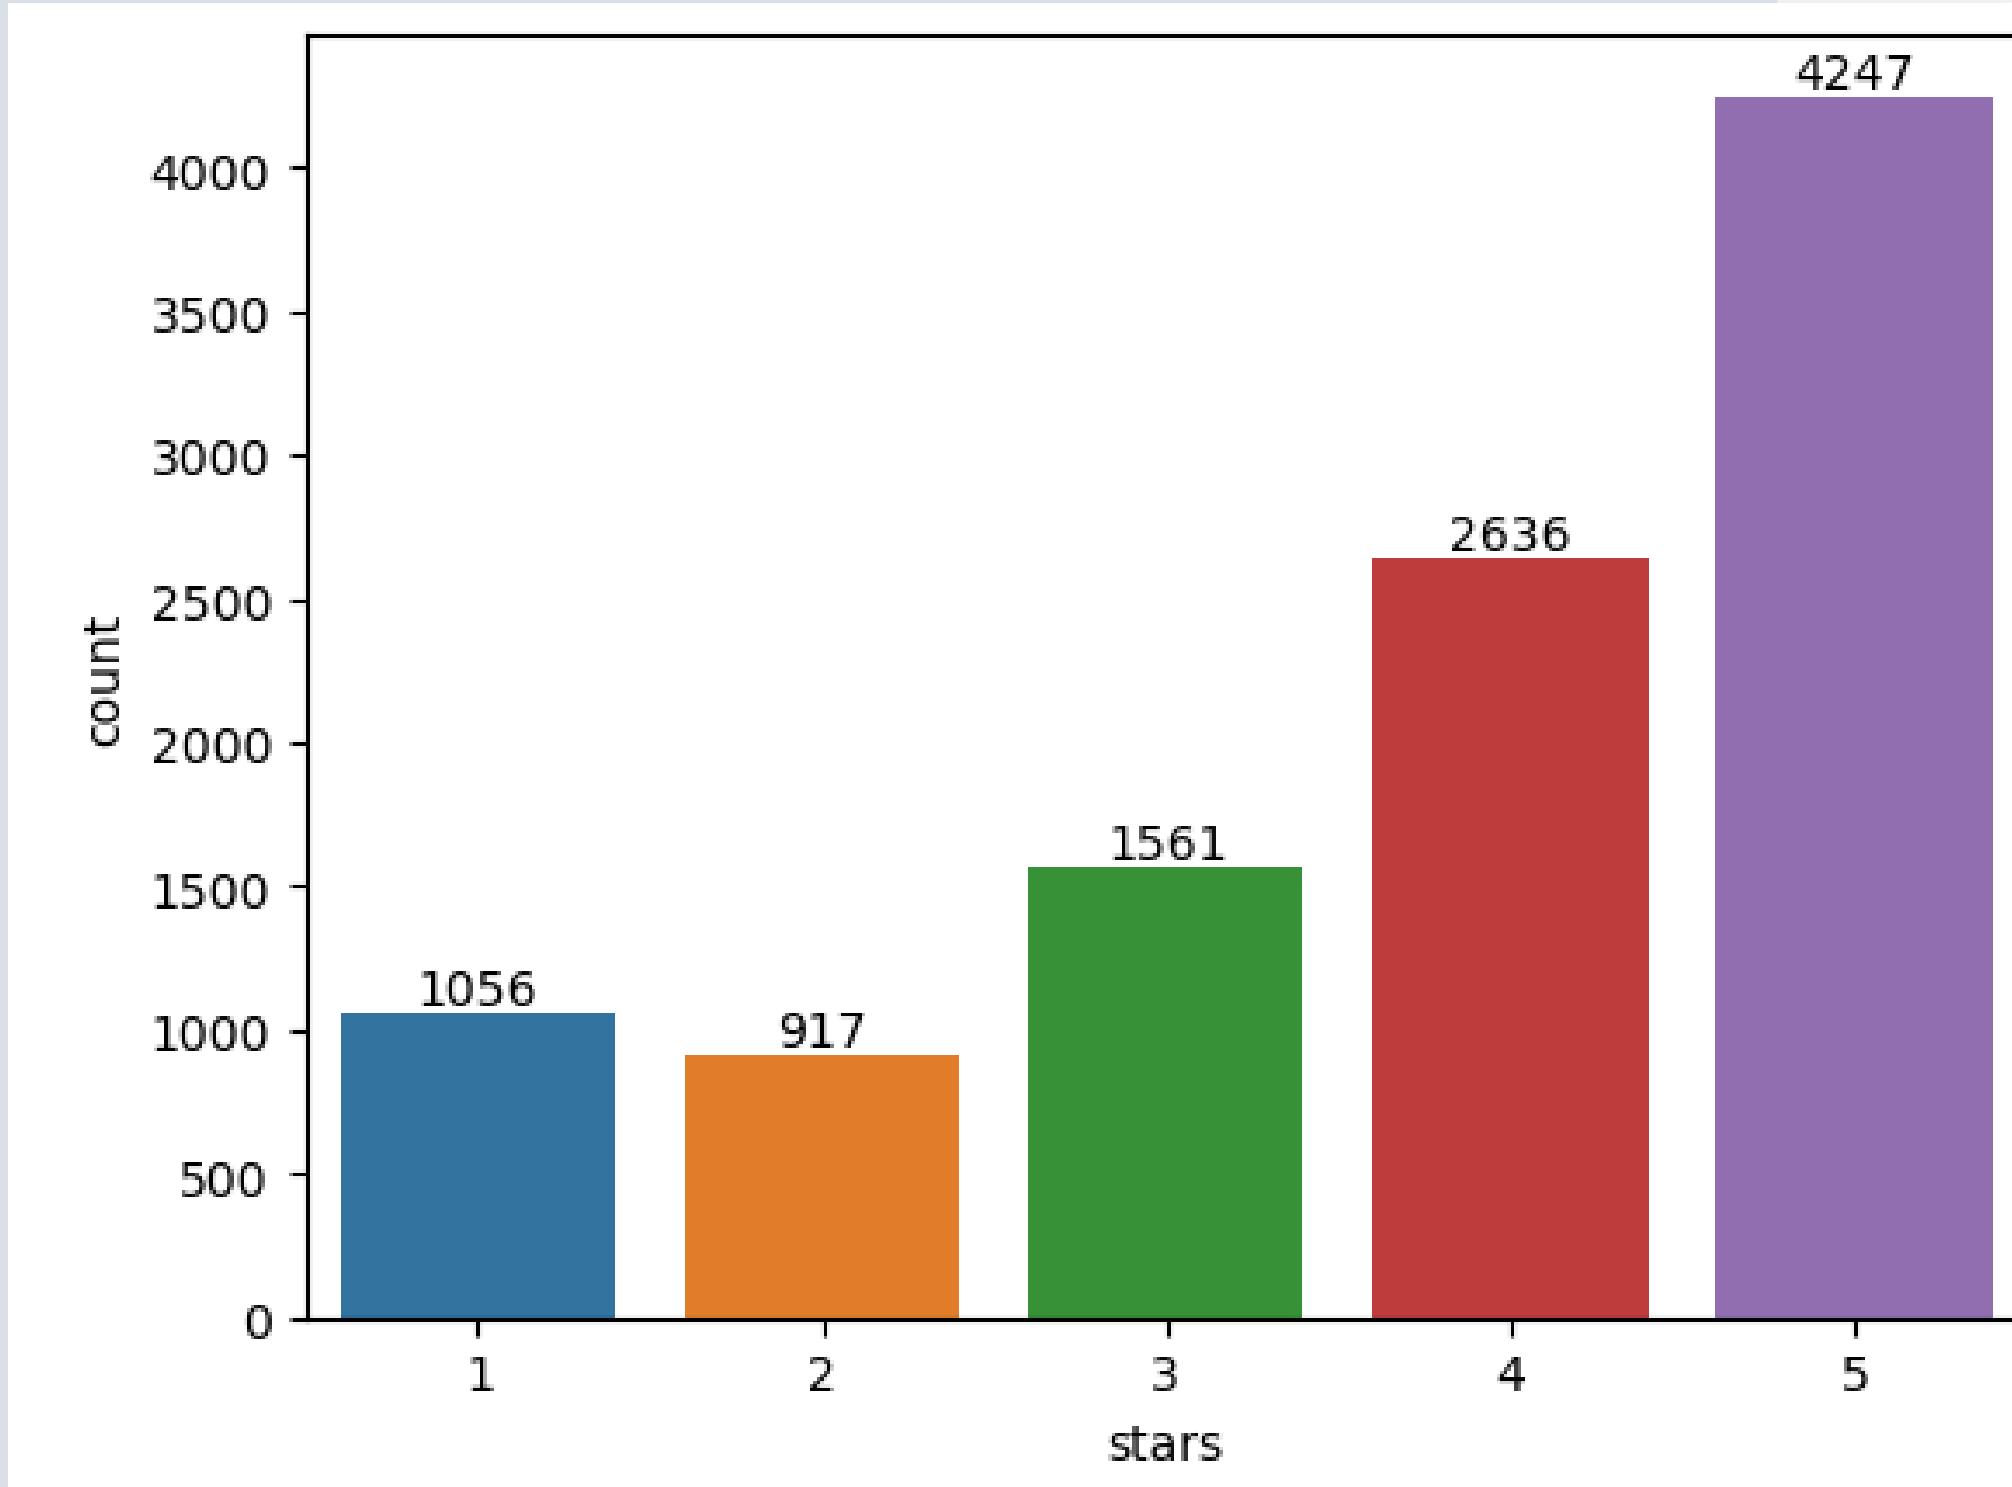
Seaborn
Data visualisations





Results

After dropping unnecessary columns

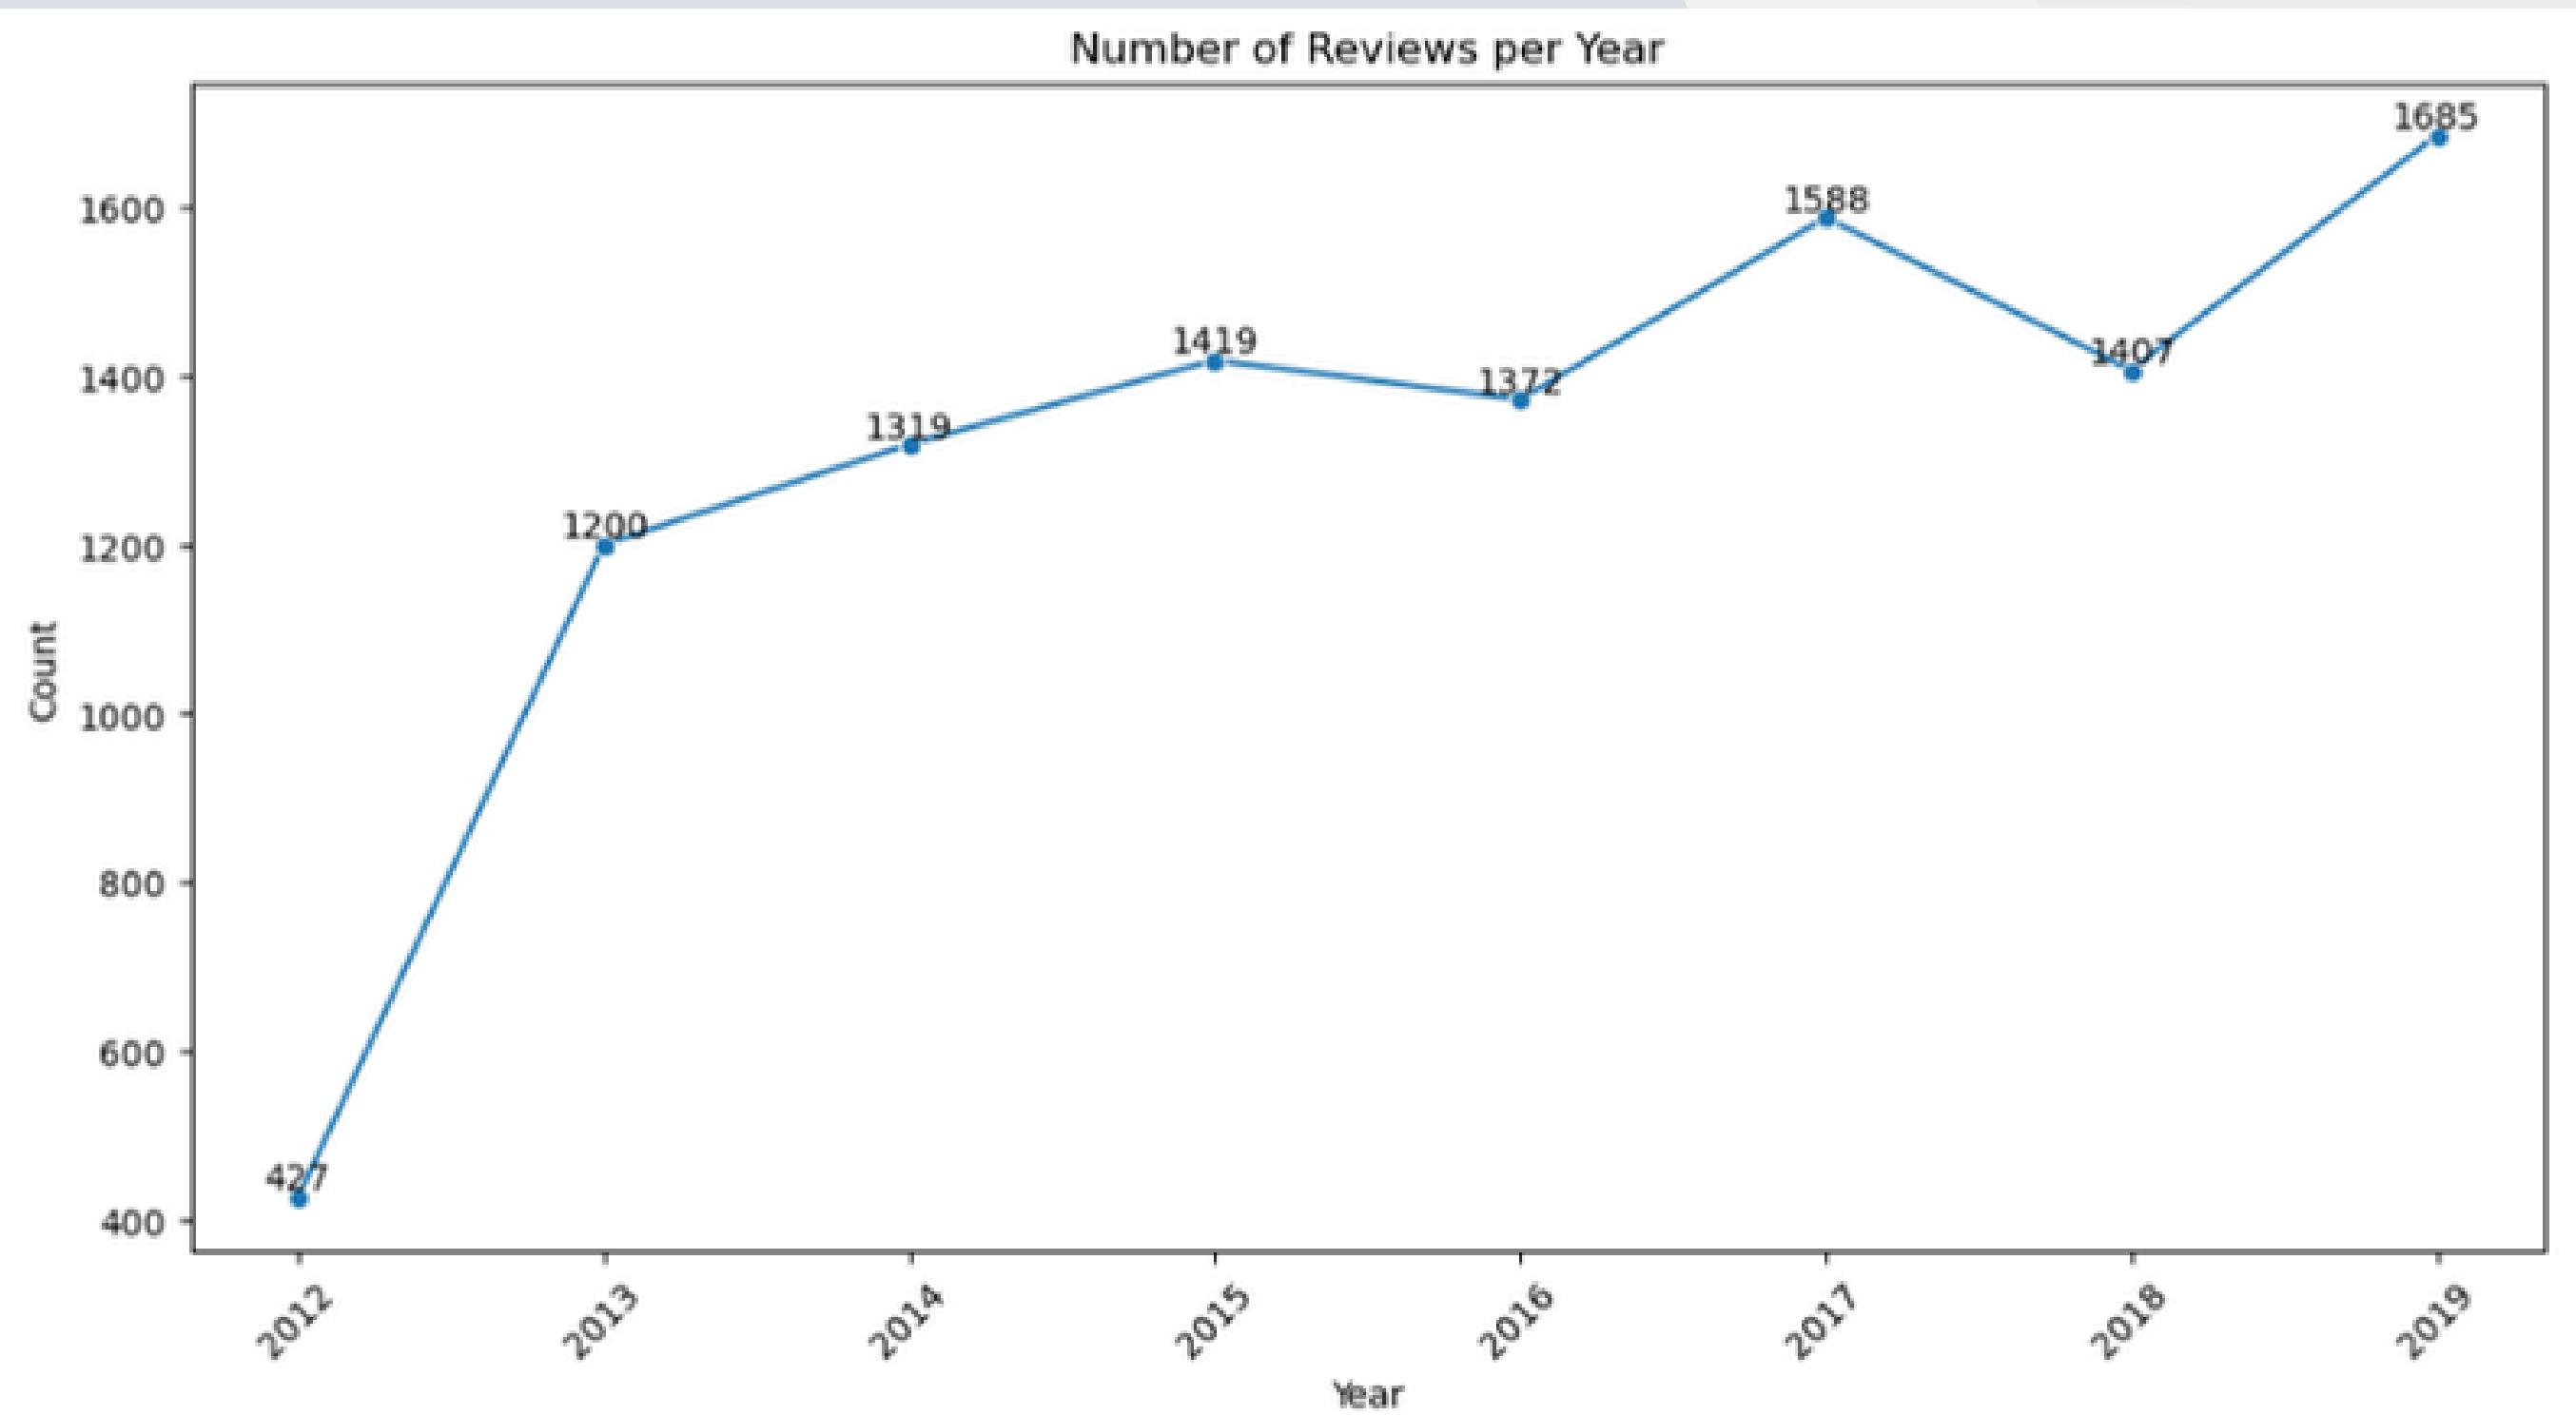


Exploratory
Data Analysis

Stars distribution
before text processing



Results

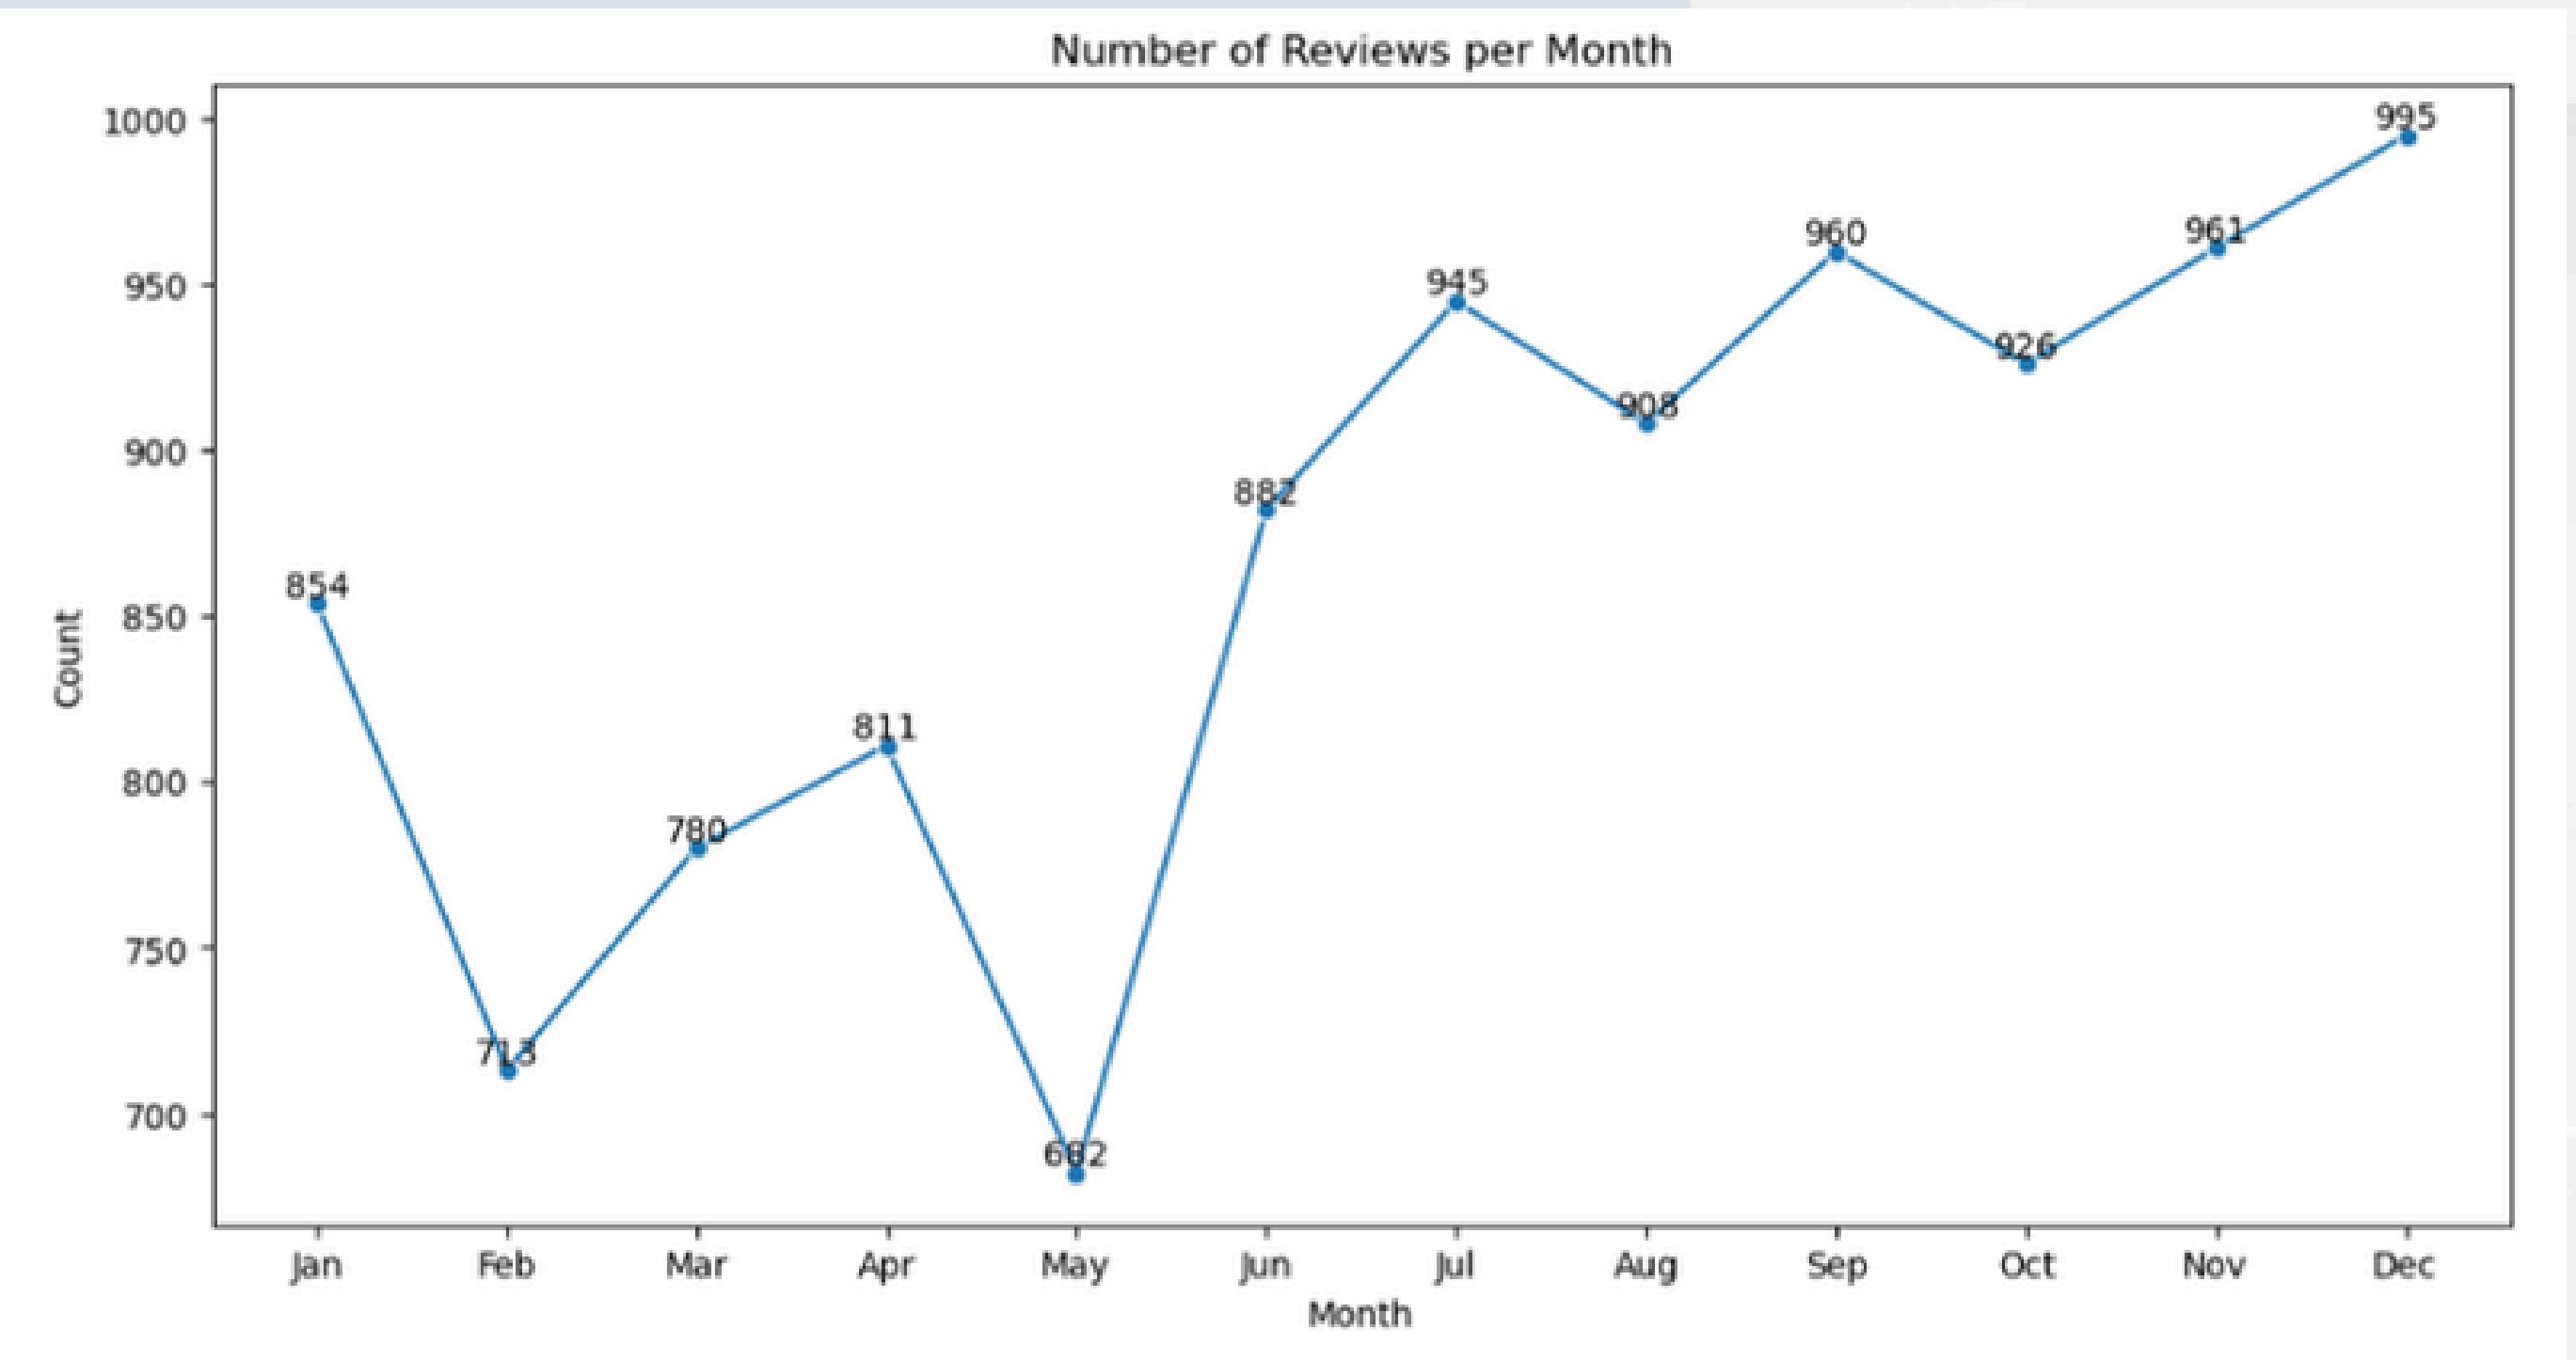


Exploratory
Data Analysis

Number of Reviews
per Year



Results

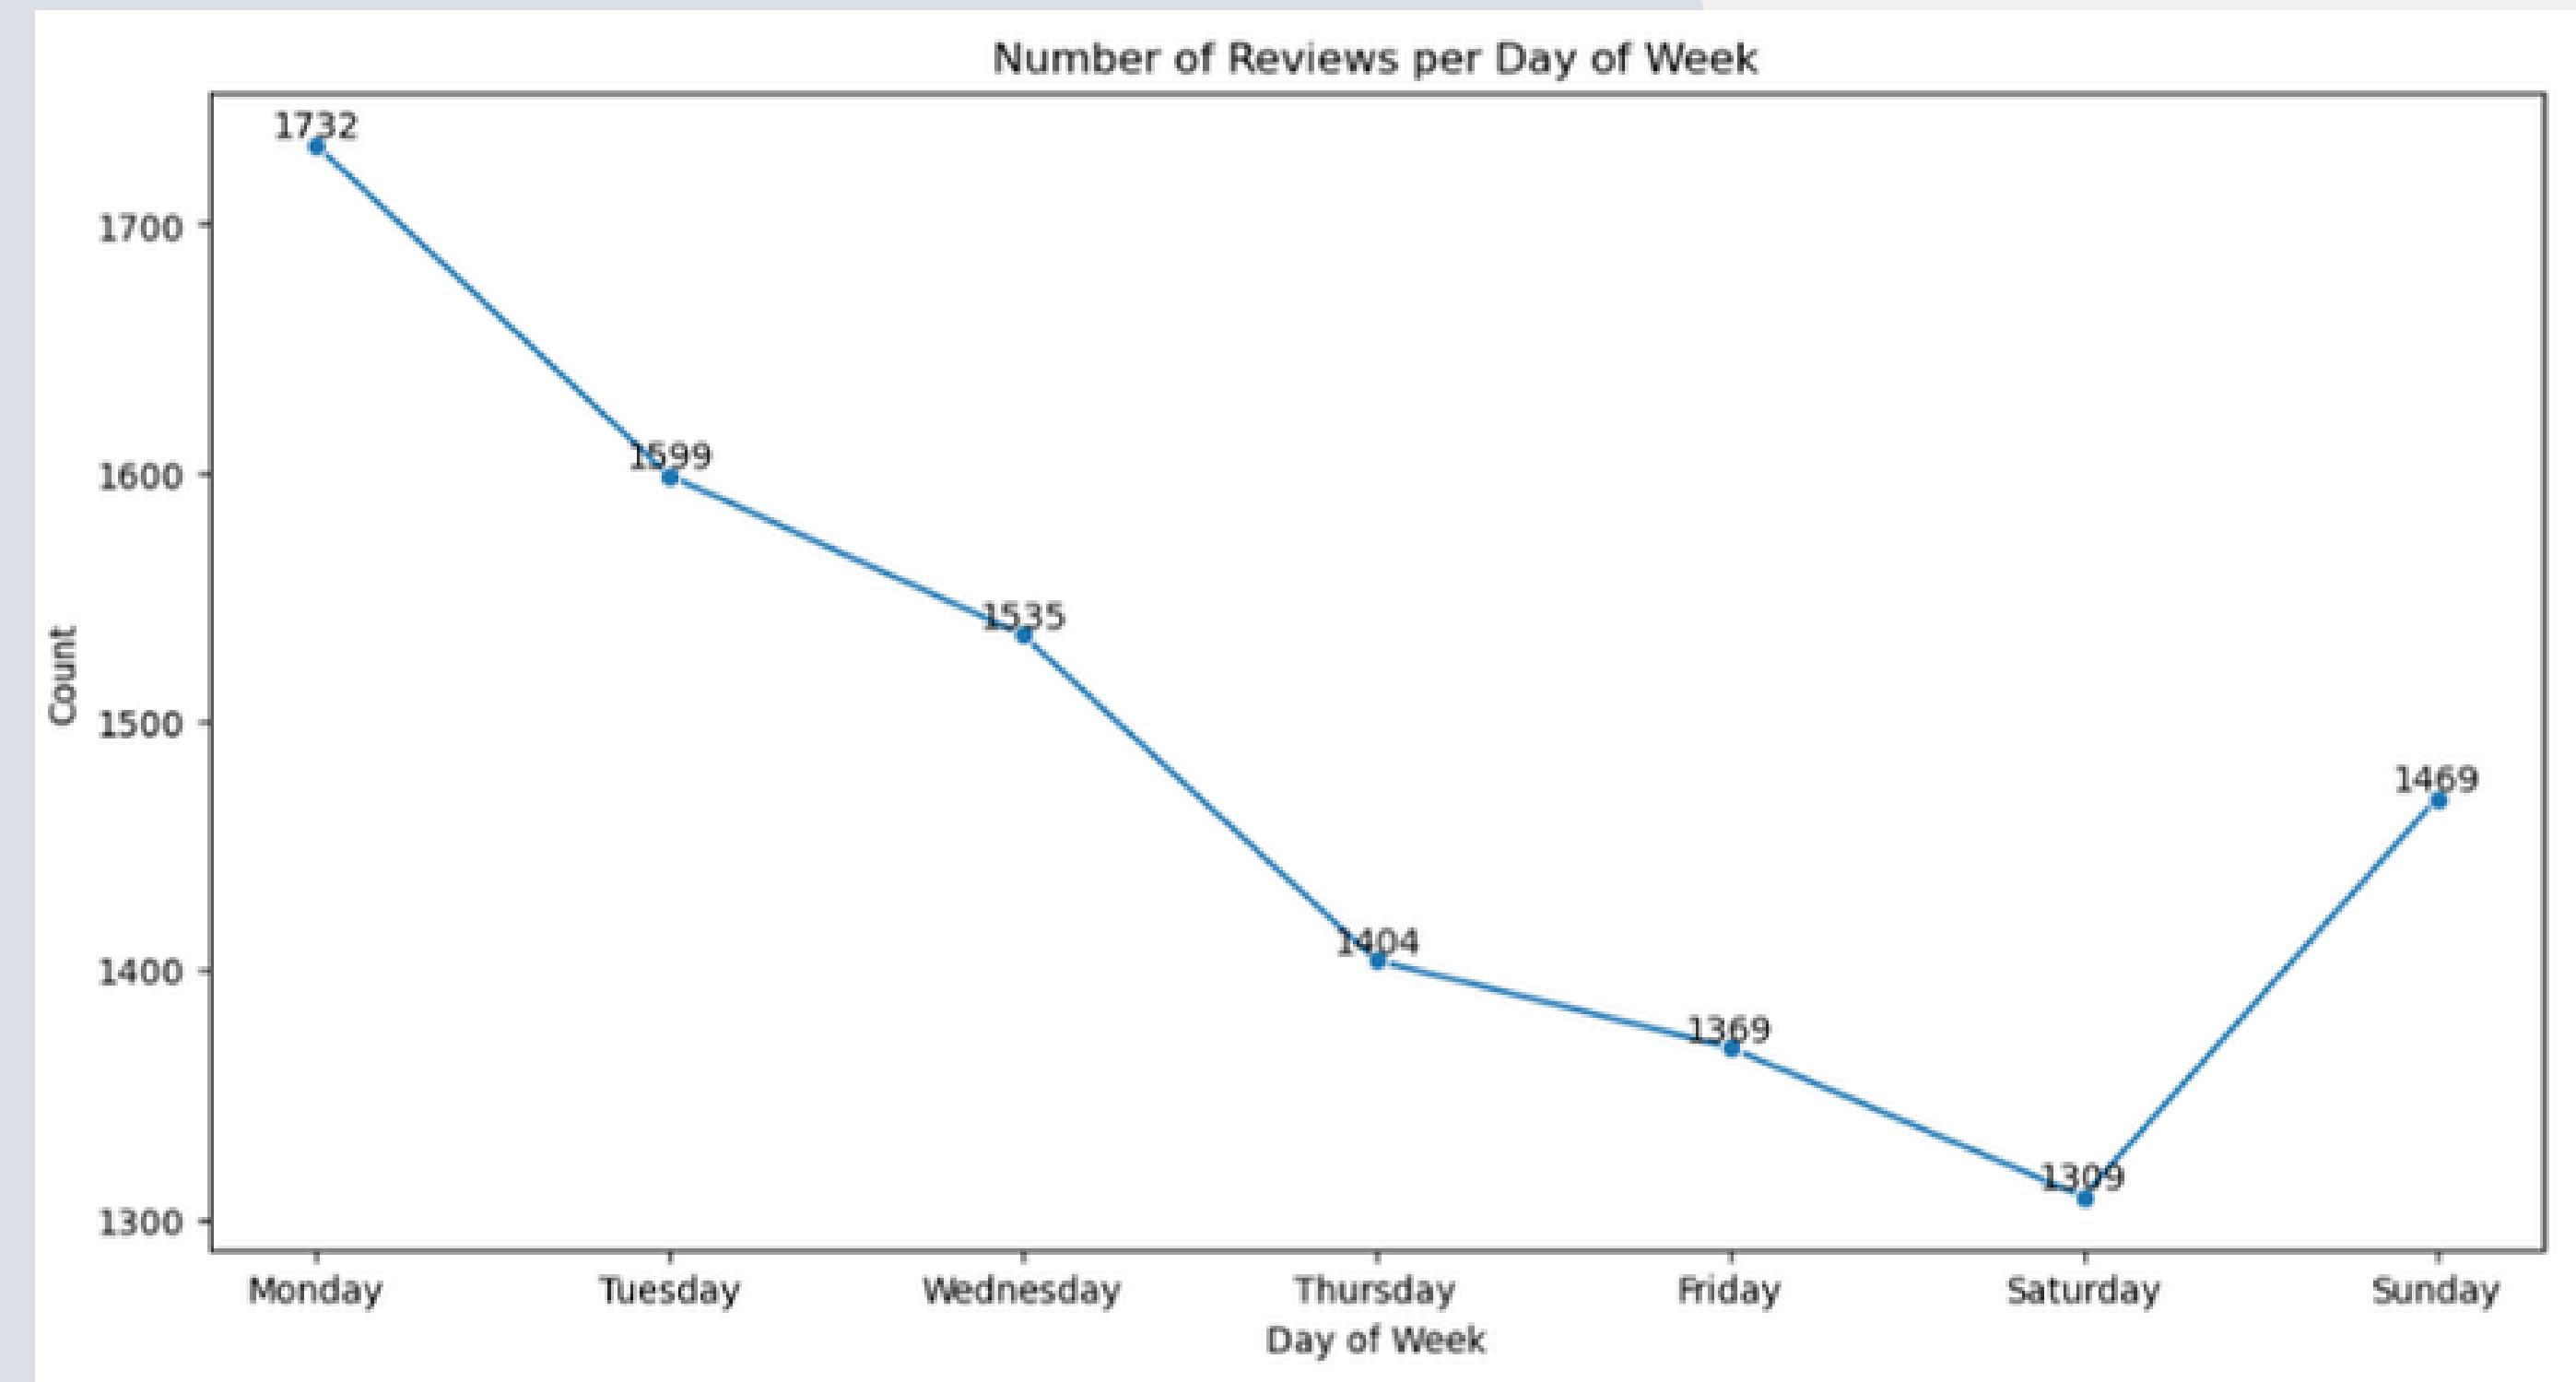


Exploratory
Data Analysis

Number of Reviews
per Month



Results

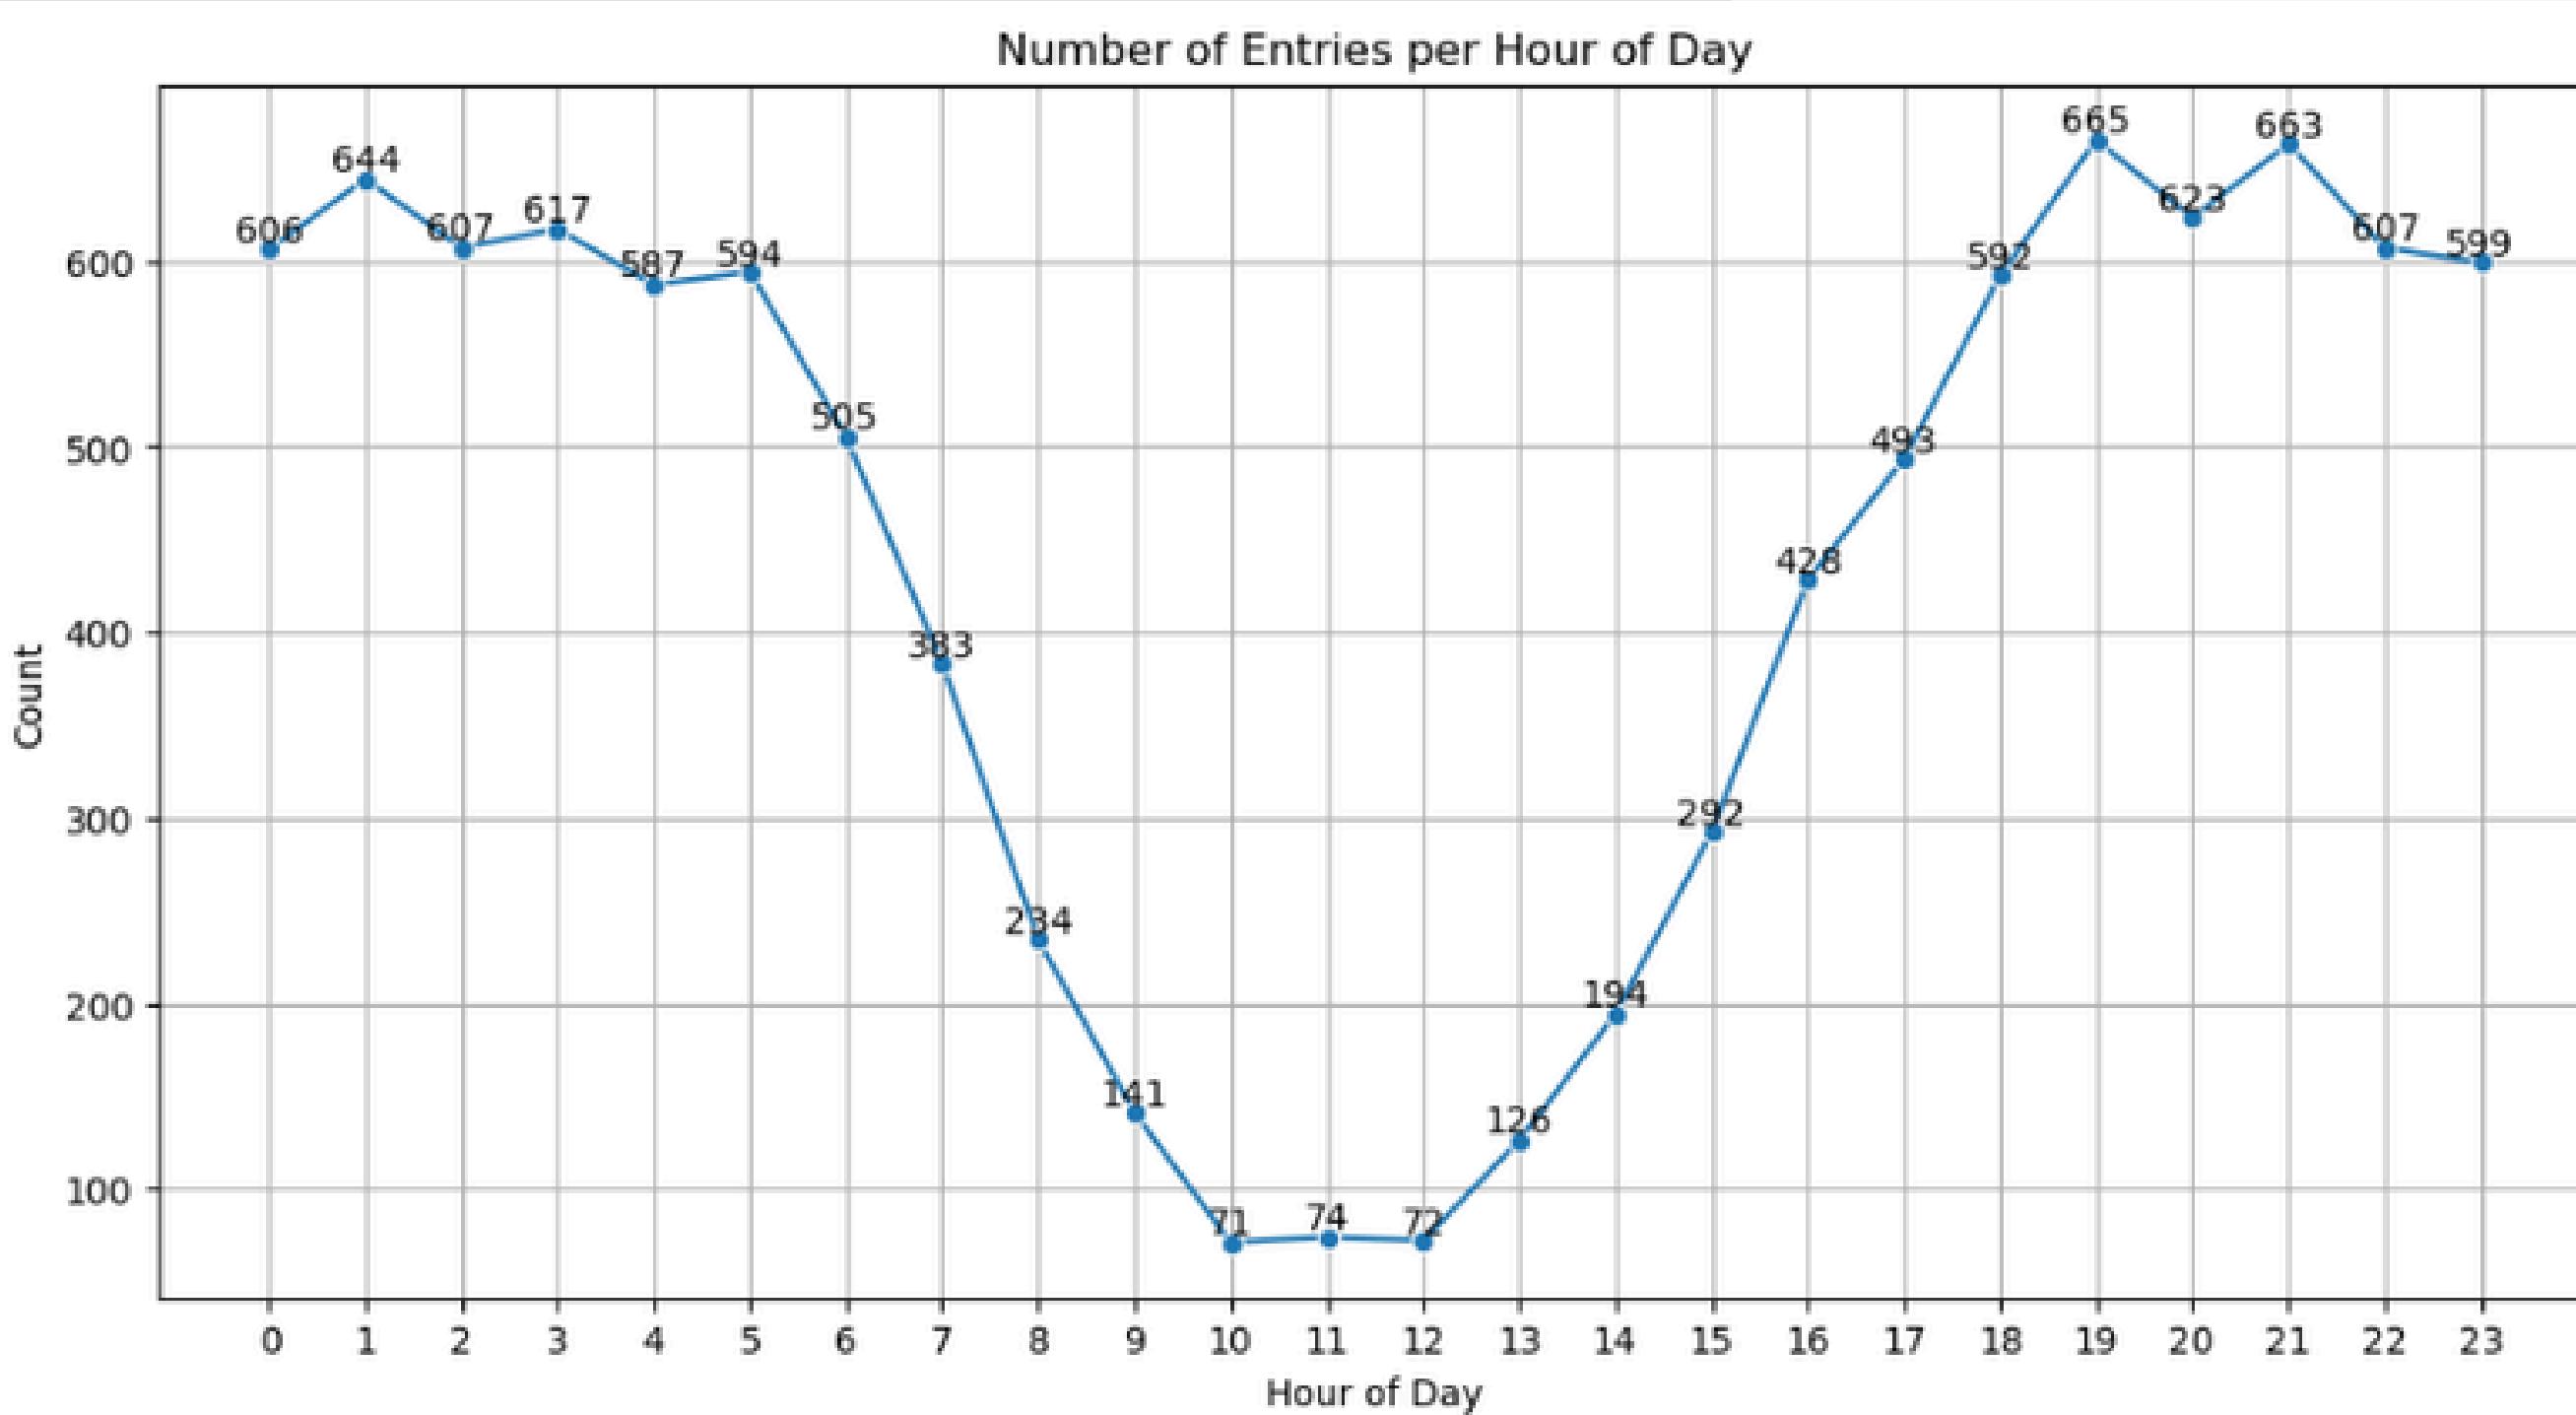


Exploratory
Data Analysis

Number of Reviews
per DayofWeek



Results

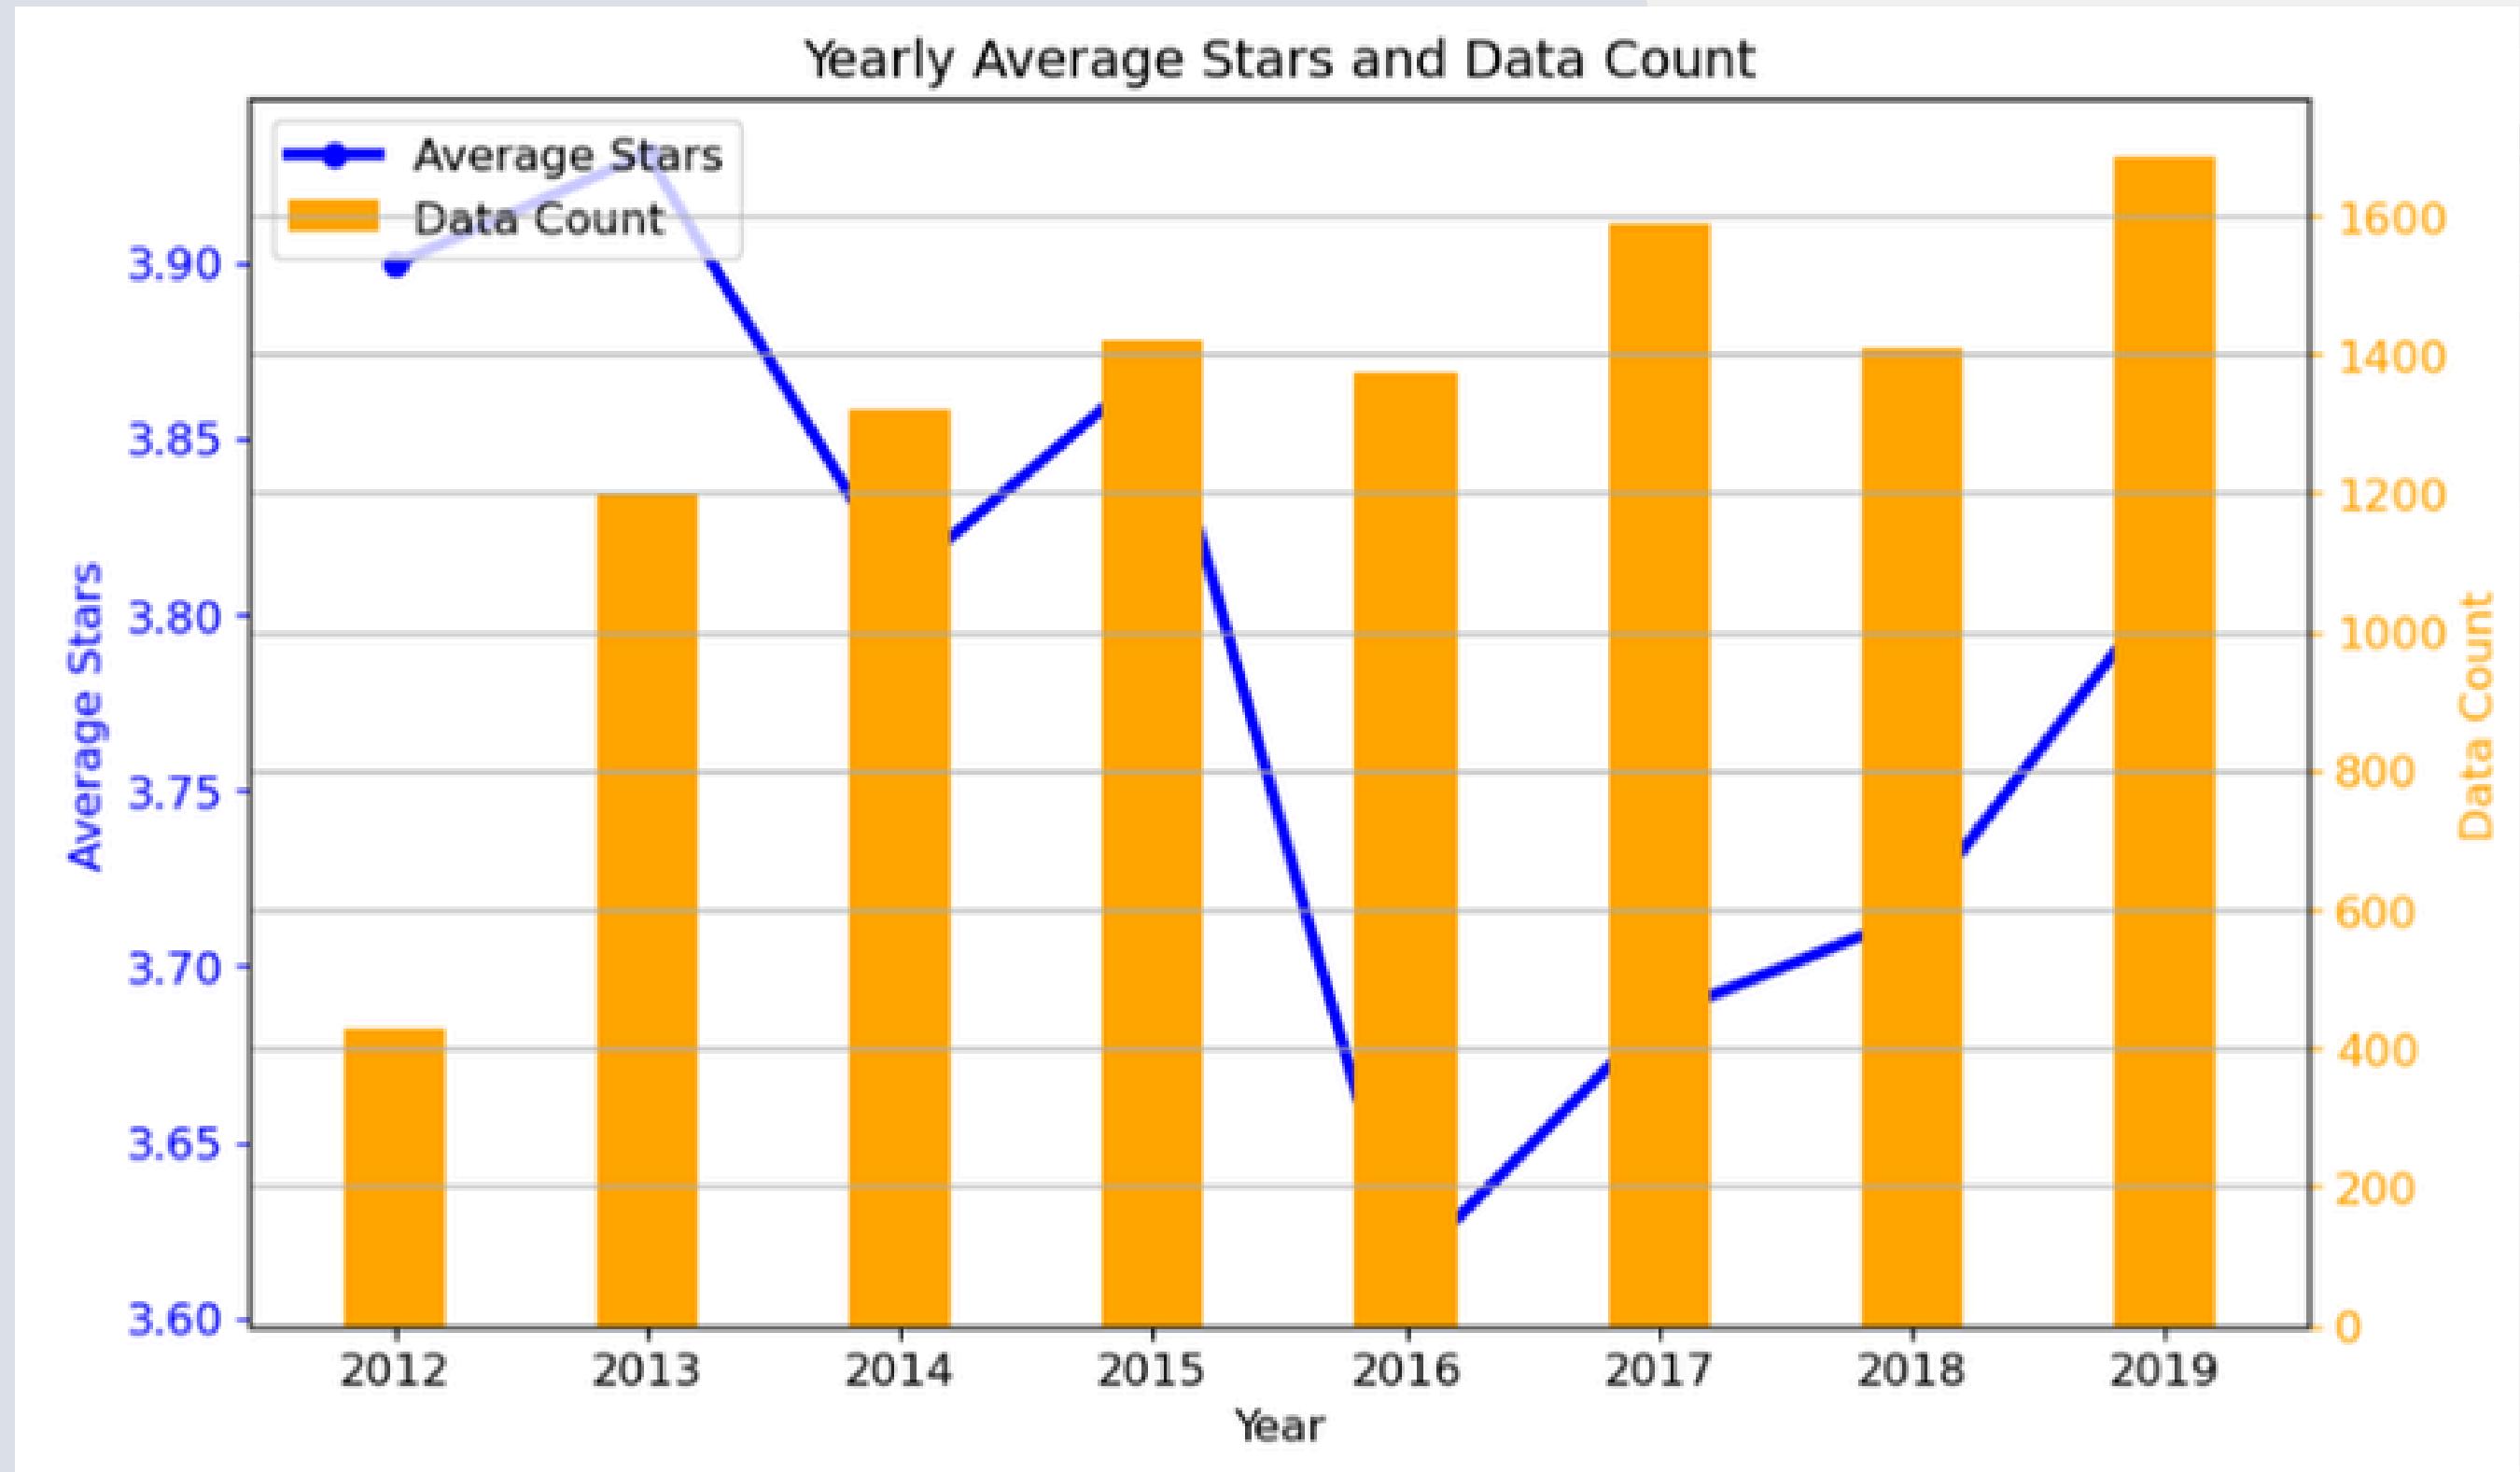


Exploratory
Data Analysis

Number of Reviews
along a Day



Results



Exploratory
Data Analysis

Yearly avg. stars
and Data Count

Results

Stemming & Lemmatisation

Detecting Language

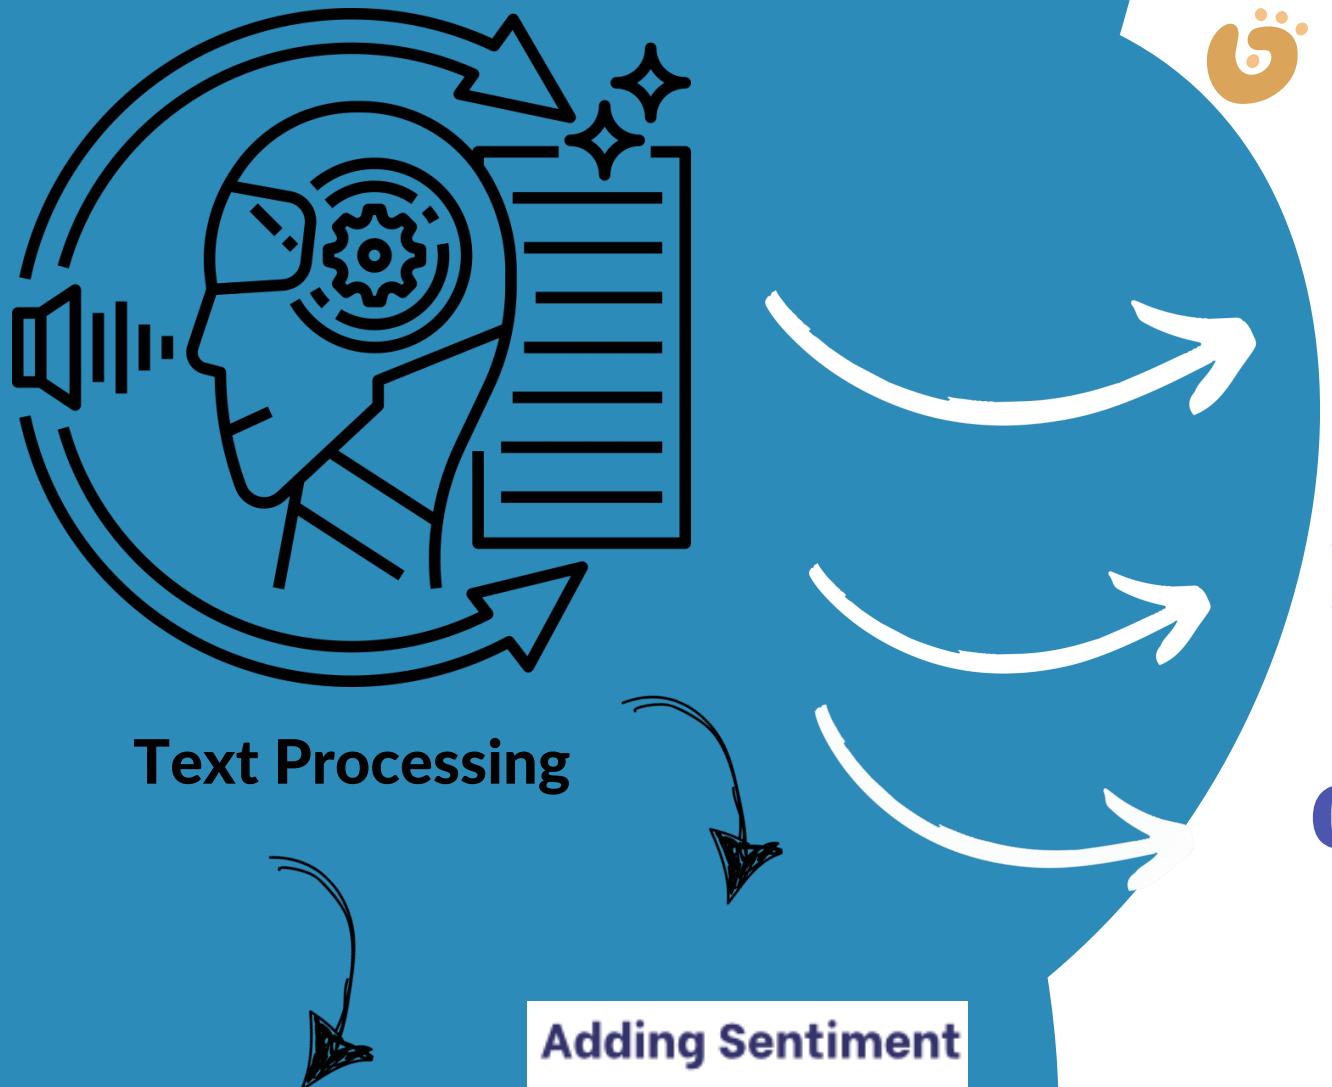
Convert Lower Case

Removing Characters

- numbers (1-9)
- punctuation ()
- new line (\n)

Adding Sentiment

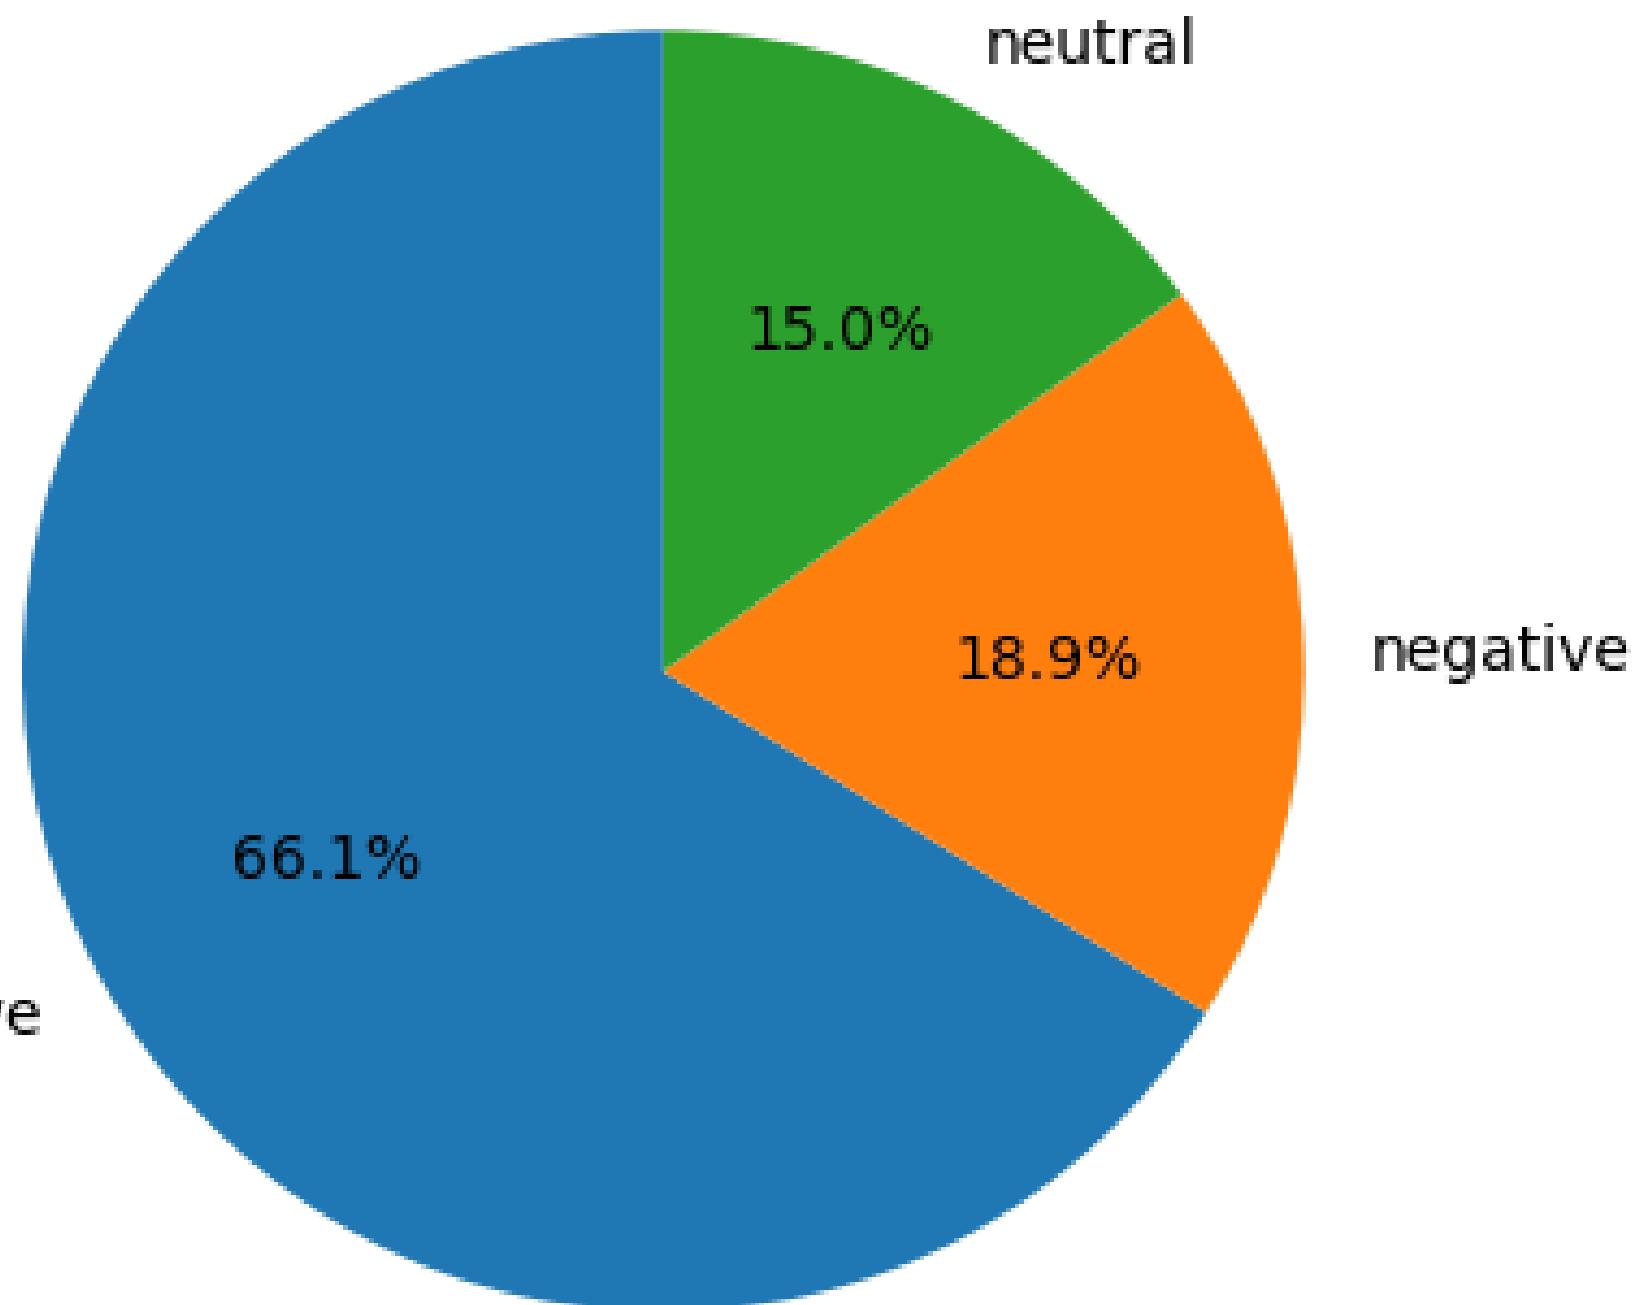
- 1,2 → Negative
- 3 → Neutral
- 4,5 → Positive





Results

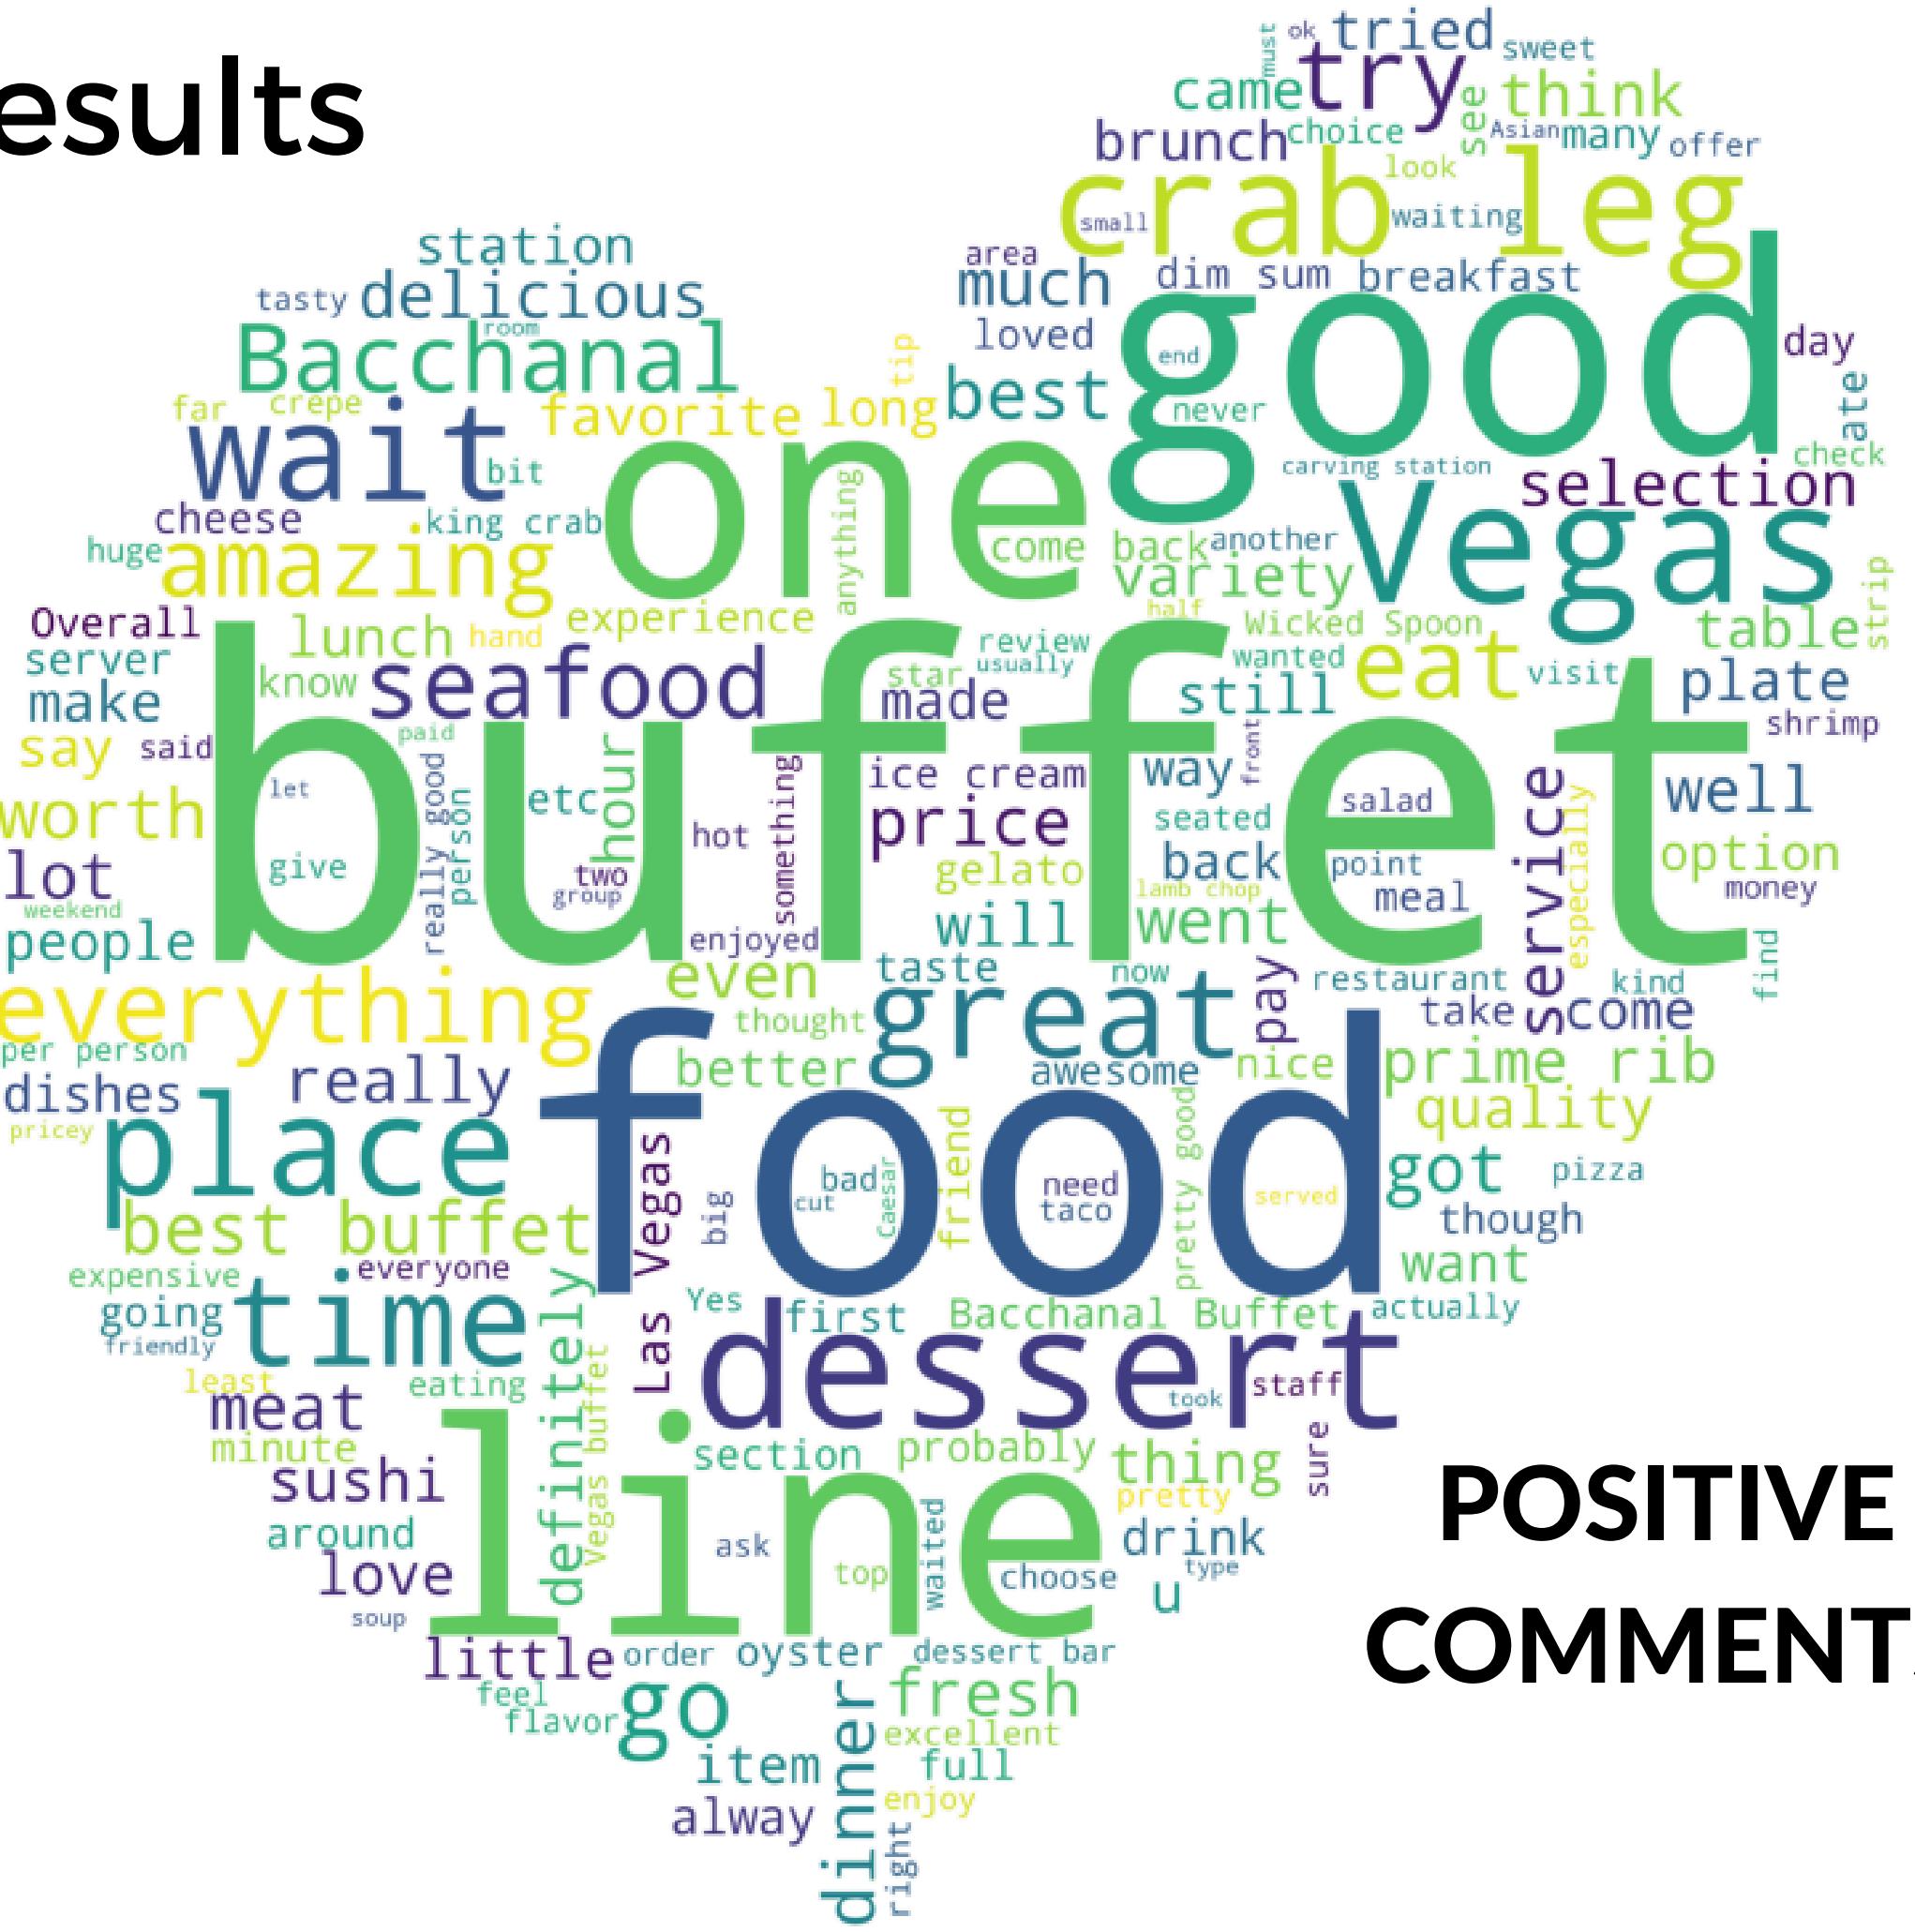
Sentiment Distribution



Sentiment distribution
after text processing

Results

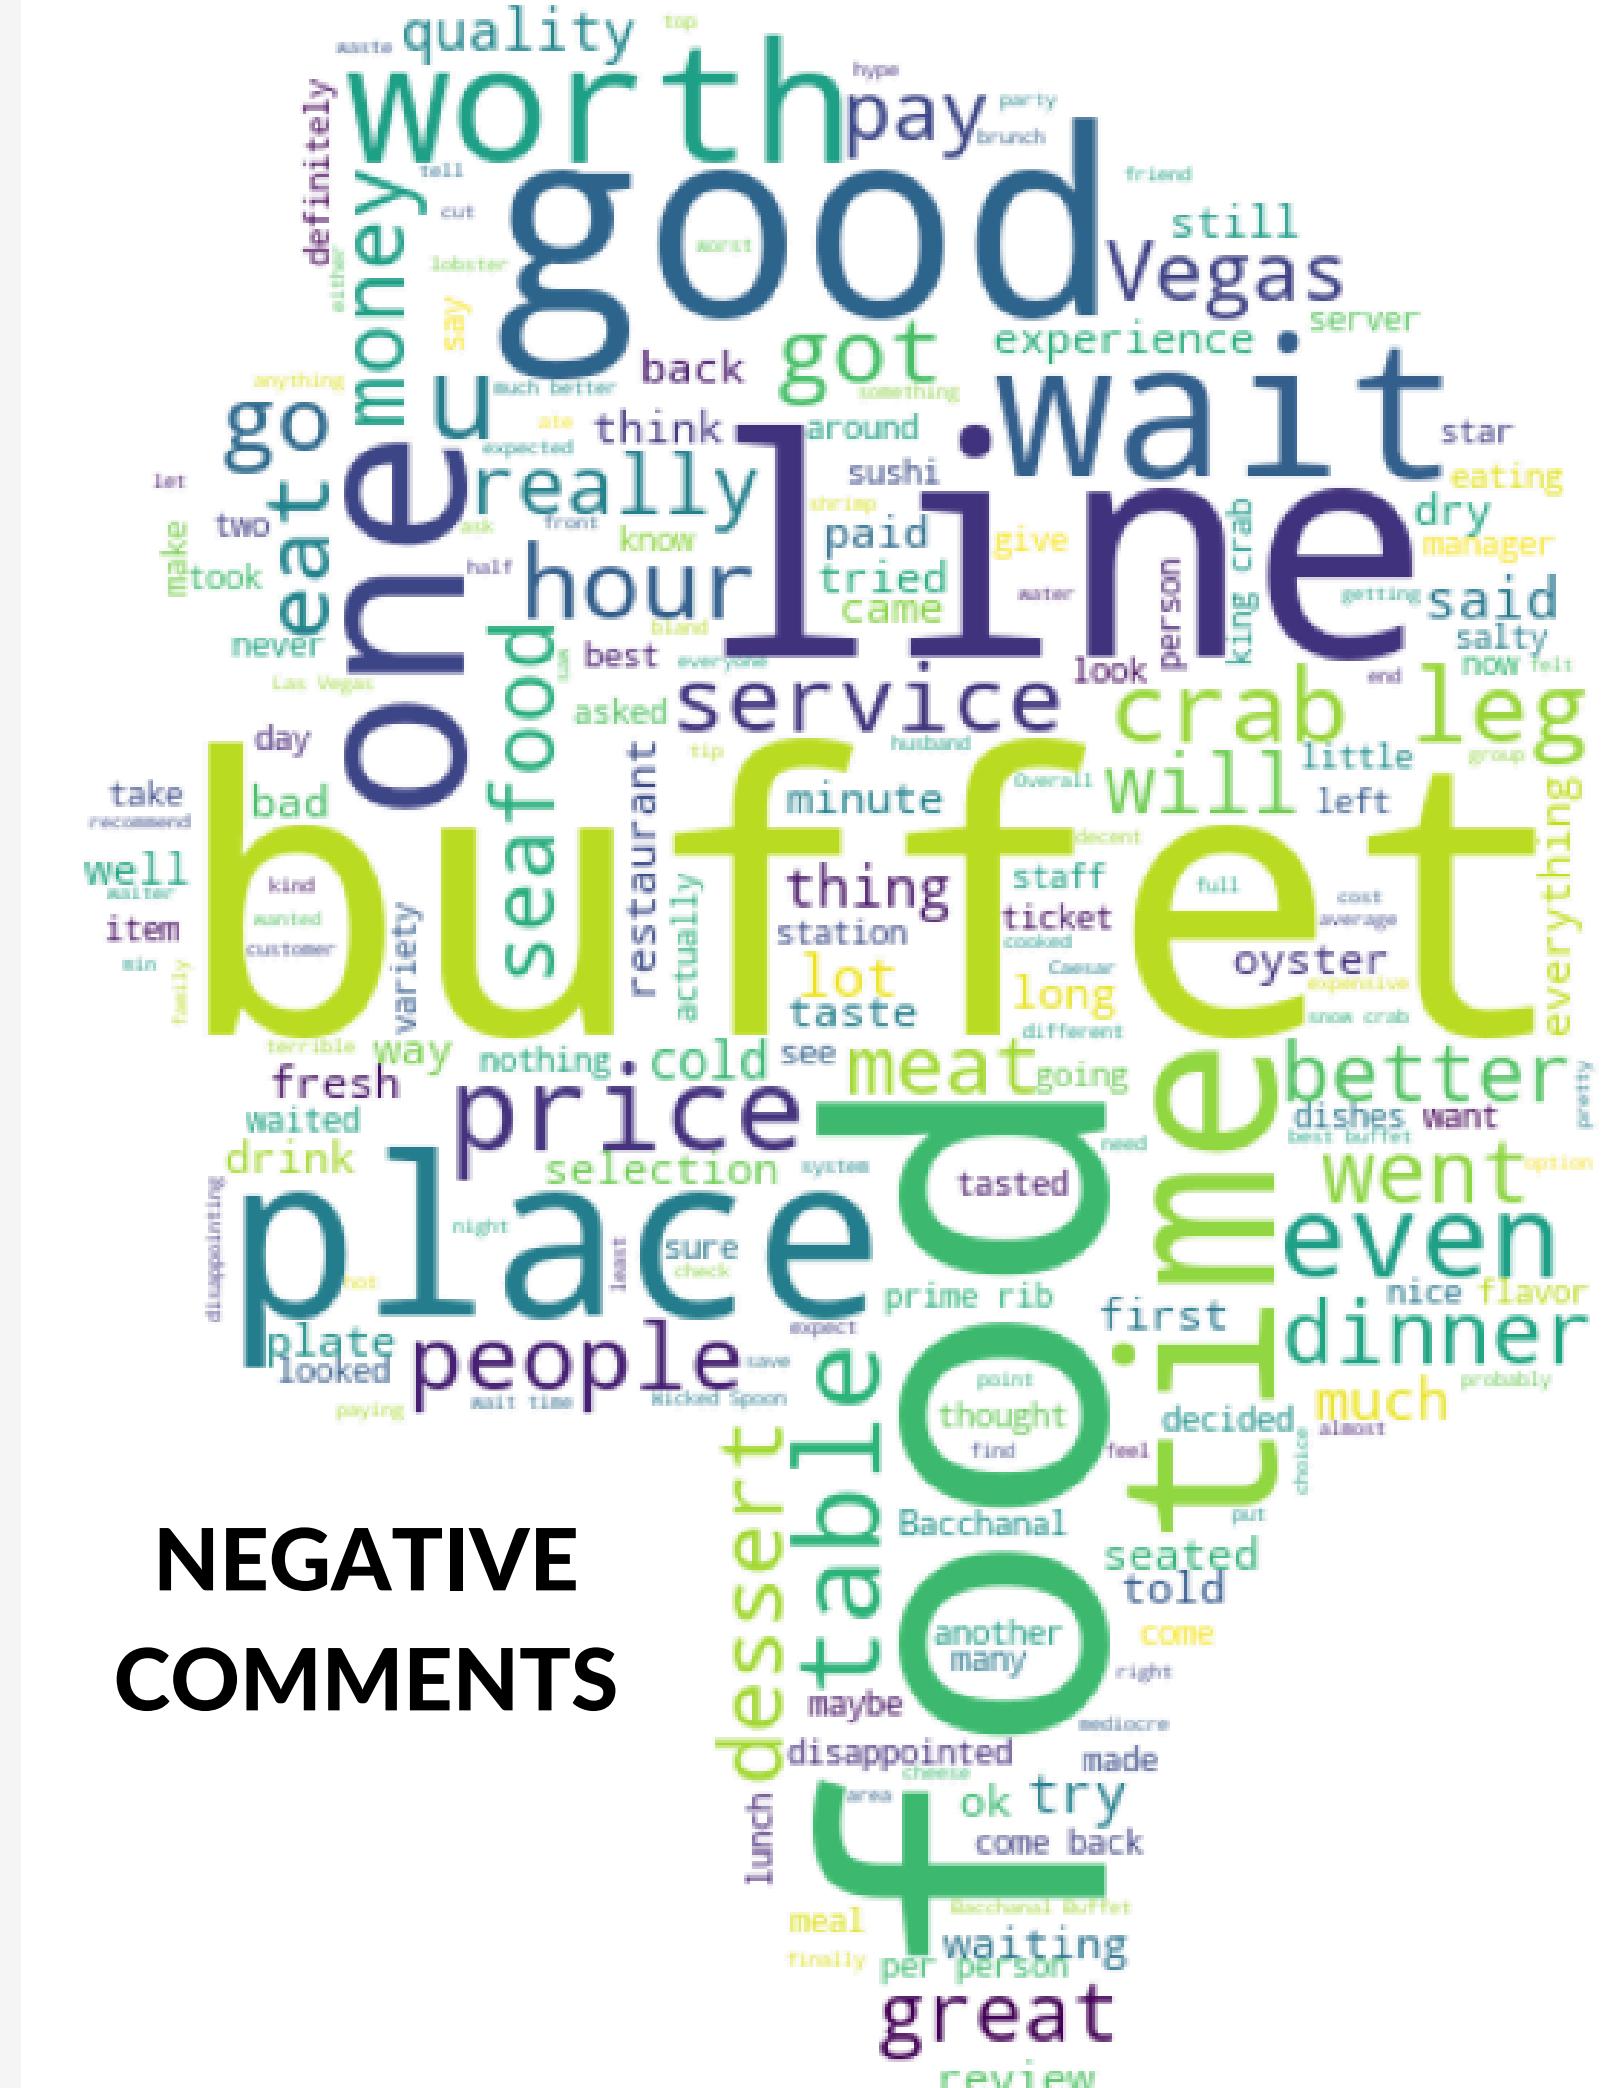
5-STAR COMMENTS



Results



1-STAR COMMENTS





Results

`LogisticRegression()`,
`KNeighborsClassifier()`,
`DecisionTreeClassifier()`,
`RandomForestClassifier()`,
`AdaBoostClassifier()`,
`GradientBoostingClassifier()`,
`MultinomialNB()`,
`BernoulliNB()`

Methods

CountVectorizer with unigram
CountVectorizer with unigram and lemmafn
CountVectorizer with bigram
CountVectorizer with bigram and lemmafn
TfidVectorizer with unigram
TfidVectorizer with unigram and lemmafn
TfidVectorizer with bigram
TfidVectorizer with bigram and lemmafn

Algorithms

Results

CountVectorizes with/out ngrams or lemma

```
# 1st method: CountVectorizer with unigram  
find_best_model(x, y)
```

```
('Logistic Regression',  
 {'accuracy': 0.8072289156626506, ←  
  'f1_score': 0.8004543680700585,  
  'confusion_matrix': array([[ 298,    86,    30],  
     [ 74,   128,    60],  
     [ 43,   107, 1249]], dtype=int64),  
  'classification_report': 'precision    recall  f1-score  
0.72      0.72      0.72      414\n      neutral      0.40      0.42  
0.93      0.89      0.91     1399\n      accuracy  
g       0.68      0.70      0.69      2075\nweighted avg      0.82
```

```
# 2nd method: CountVectorizer with unigram and Lemmatfn  
find_best_model(xs, y)
```

```
('Logistic Regression',  
 {'accuracy': 0.7850602409638554,  
  'f1_score': 0.7764937243114267,  
  'confusion_matrix': array([[ 278,    81,    29],  
     [ 82,   111,    70],  
     [ 55,   129, 1240]], dtype=int64),  
  'classification_report': 'precision    recall  f1-score  
0.67      0.72      0.69      388\n      neutral      0.35      0.42  
0.93      0.87      0.90     1424\n      accuracy  
g       0.65      0.67      0.66      2075\nweighted avg      0.80
```

```
# 3rd method: CountVectorizer with bigram  
find_best_model(x, y)  
  
('Logistic Regression',  
 {'accuracy': 0.7966265060240963,  
  'f1_score': 0.788418047812878,  
  'confusion_matrix': array([[ 285,    80,    24],  
     [ 75,   120,    67],  
     [ 55,   121, 1248]], dtype=int64),  
  'classification_report': 'precision    recall  f1-score  
0.69      0.73      0.71      389\n      neutral      0.37      0.46  
0.93      0.88      0.90     1424\n      accuracy  
g       0.66      0.69      0.67      2075\nweighted avg      0.82
```

```
# 4th method: CountVectorizer with bigram and Lemmafn  
find_best_model(xs, y)  
  
('Logistic Regression',  
 {'accuracy': 0.792289156626506,  
  'f1_score': 0.7834579175763278,  
  'confusion_matrix': array([[ 281,    81,    27],  
     [ 80,   115,    64],  
     [ 54,   125, 1248]], dtype=int64),  
  'classification_report': 'precision    recall  f1-score  
0.68      0.72      0.70      389\n      neutral      0.36      0.44  
0.93      0.87      0.90     1427\n      accuracy  
g       0.66      0.68      0.67      2075\nweighted avg      0.81
```

6 Results

TfidfVectorizes with/out ngrams or lemma

```
# 5th method: TfidfVectorizer with unigram
find_best_model(x, y)

('Logistic Regression',
 {'accuracy': 0.8144578313253013,
  'f1_score': 0.7935681003686282,
  'confusion_matrix': array([[ 295,    68,    16],
                             [ 41,   96,   24],
                             [ 79,  157, 1299]], dtype=int64),
  'classification_report': 'precision      recall
                           0.71      0.78      0.74      379\n      neutral     0.30
                           0.97      0.85      0.90      1535\n\n      accuracy
                           g       0.66      0.74      0.68      2075\nweighted avg'}
```

```
# 6th method: TfidfVectorizer with unigram and Lemmaf
find_best_model(xs, y)

('Logistic Regression',
 {'accuracy': 0.8028915662650602,
  'f1_score': 0.7787191133218953,
  'confusion_matrix': array([[ 287,    70,    17],
                             [ 49,   81,   24],
                             [ 79,  170, 1298]], dtype=int64),
  'classification_report': 'precision      recall
                           0.69      0.77      0.73      374\n      neutral     0.25
                           0.97      0.84      0.90      1547\n\n      accuracy
                           g       0.64      0.71      0.66      2075\nweighted avg'}
```

```
# 7th method: TfidfVectorizer with bigram
find_best_model(x, y)

('Logistic Regression',
 {'accuracy': 0.8004819277108434,
  'f1_score': 0.7730826063237212,
  'confusion_matrix': array([[ 274,    64,    10],
                             [ 48,   76,   18],
                             [ 93,  181, 1311]], dtype=int64),
  'classification_report': 'precision      recall
                           0.66      0.79      0.72      348\n      neutral     0.24
                           0.98      0.83      0.90      1585\n\n      accuracy
                           g       0.63      0.72      0.65      2075\nweighted avg'}
```

```
# 8th method: TfidfVectorizer with bigram and Lemmaf
find_best_model(xs, y)
```

```
('Logistic Regression',
 {'accuracy': 0.8004819277108434,
  'f1_score': 0.7732209243189241,
  'confusion_matrix': array([[ 285,    66,    16],
                             [ 51,   72,   19],
                             [ 79,  183, 1304]], dtype=int64),
  'classification_report': 'precision      recall
                           0.69      0.78      0.73      367\n      neutral     0.22
                           0.97      0.83      0.90      1566\n\n      accuracy
                           g       0.63      0.71      0.65      2075\nweighted avg'}
```

6 Results

```
# 5th method: TfidVectorizer with unigram
find_best_model(x, y)

('Logistic Regression',
 {'accuracy': 0.8144578313253013, ←
  'f1_score': 0.7935681003686282,
  'confusion_matrix': array([[ 295,    68,    16],
     [  41,   96,   24],
     [  79,  157, 1299]], dtype=int64),
  'classification_report': '
    precision    recall
0.71      0.78      0.74      379\n      neutral      0.30
0.97      0.85      0.90      1535\n\n      accuracy
g        0.66      0.74      0.68      2075\nweighted avg
```

```
# 6th method: TfidVectorizer with unigram and Lemmaf
find_best_model(xs, y)
```

```
('Logistic Regression',
 {'accuracy': 0.8028915662650602,
  'f1_score': 0.7787191133218953,
  'confusion_matrix': array([[ 287,    70,    17],
     [  49,   81,   24],
     [  79,  170, 1298]], dtype=int64),
  'classification_report': '
    precision    recall
0.69      0.77      0.73      374\n      neutral      0.25
0.97      0.84      0.90      1547\n\n      accuracy
g        0.64      0.71      0.66      2075\nweighted avg
```

As a result, TfidVect. with unigram provides best accuracy

Conclusion

POSITIVES

UNIGRAMS

WORD COUNTS

buffet	391.910317
food	342.609191
good	261.634207
vegas	247.394345
best	238.342973
great	221.828242
worth	194.307755
wait	194.115807
crab	192.557076
time	187.840749

BIGRAMS

WORD COUNTS

crab legs	86.314311
best buffet	85.819733
buffet vegas	56.919096
las vegas	47.284310
prime rib	46.548630
bacchanal buffet	31.381918
wicked spoon	30.290632
quality food	29.182357
dim sum	29.101293
king crab	27.711923

TRIGRAMS

WORD COUNTS

best buffet vegas	34.635155
buffet las vegas	21.661086
best buffet ve	18.535559
king crab legs	17.521715
best buffet las	15.893549
hands best buffet	11.113048
favorite buffet vegas	10.383306
best buffet strip	10.248119
far best buffet	9.493477
snow crab legs	9.266184

Conclusion

NEGATIVES

UNIGRAMS

WORD COUNTS

food	107.979998
buffet	98.069376
line	74.342726
good	68.722153
wait	67.332366
time	58.486110
crab	57.632429
just	54.318752
worth	51.350148
place	51.292444

BIGRAMS

WORD COUNTS

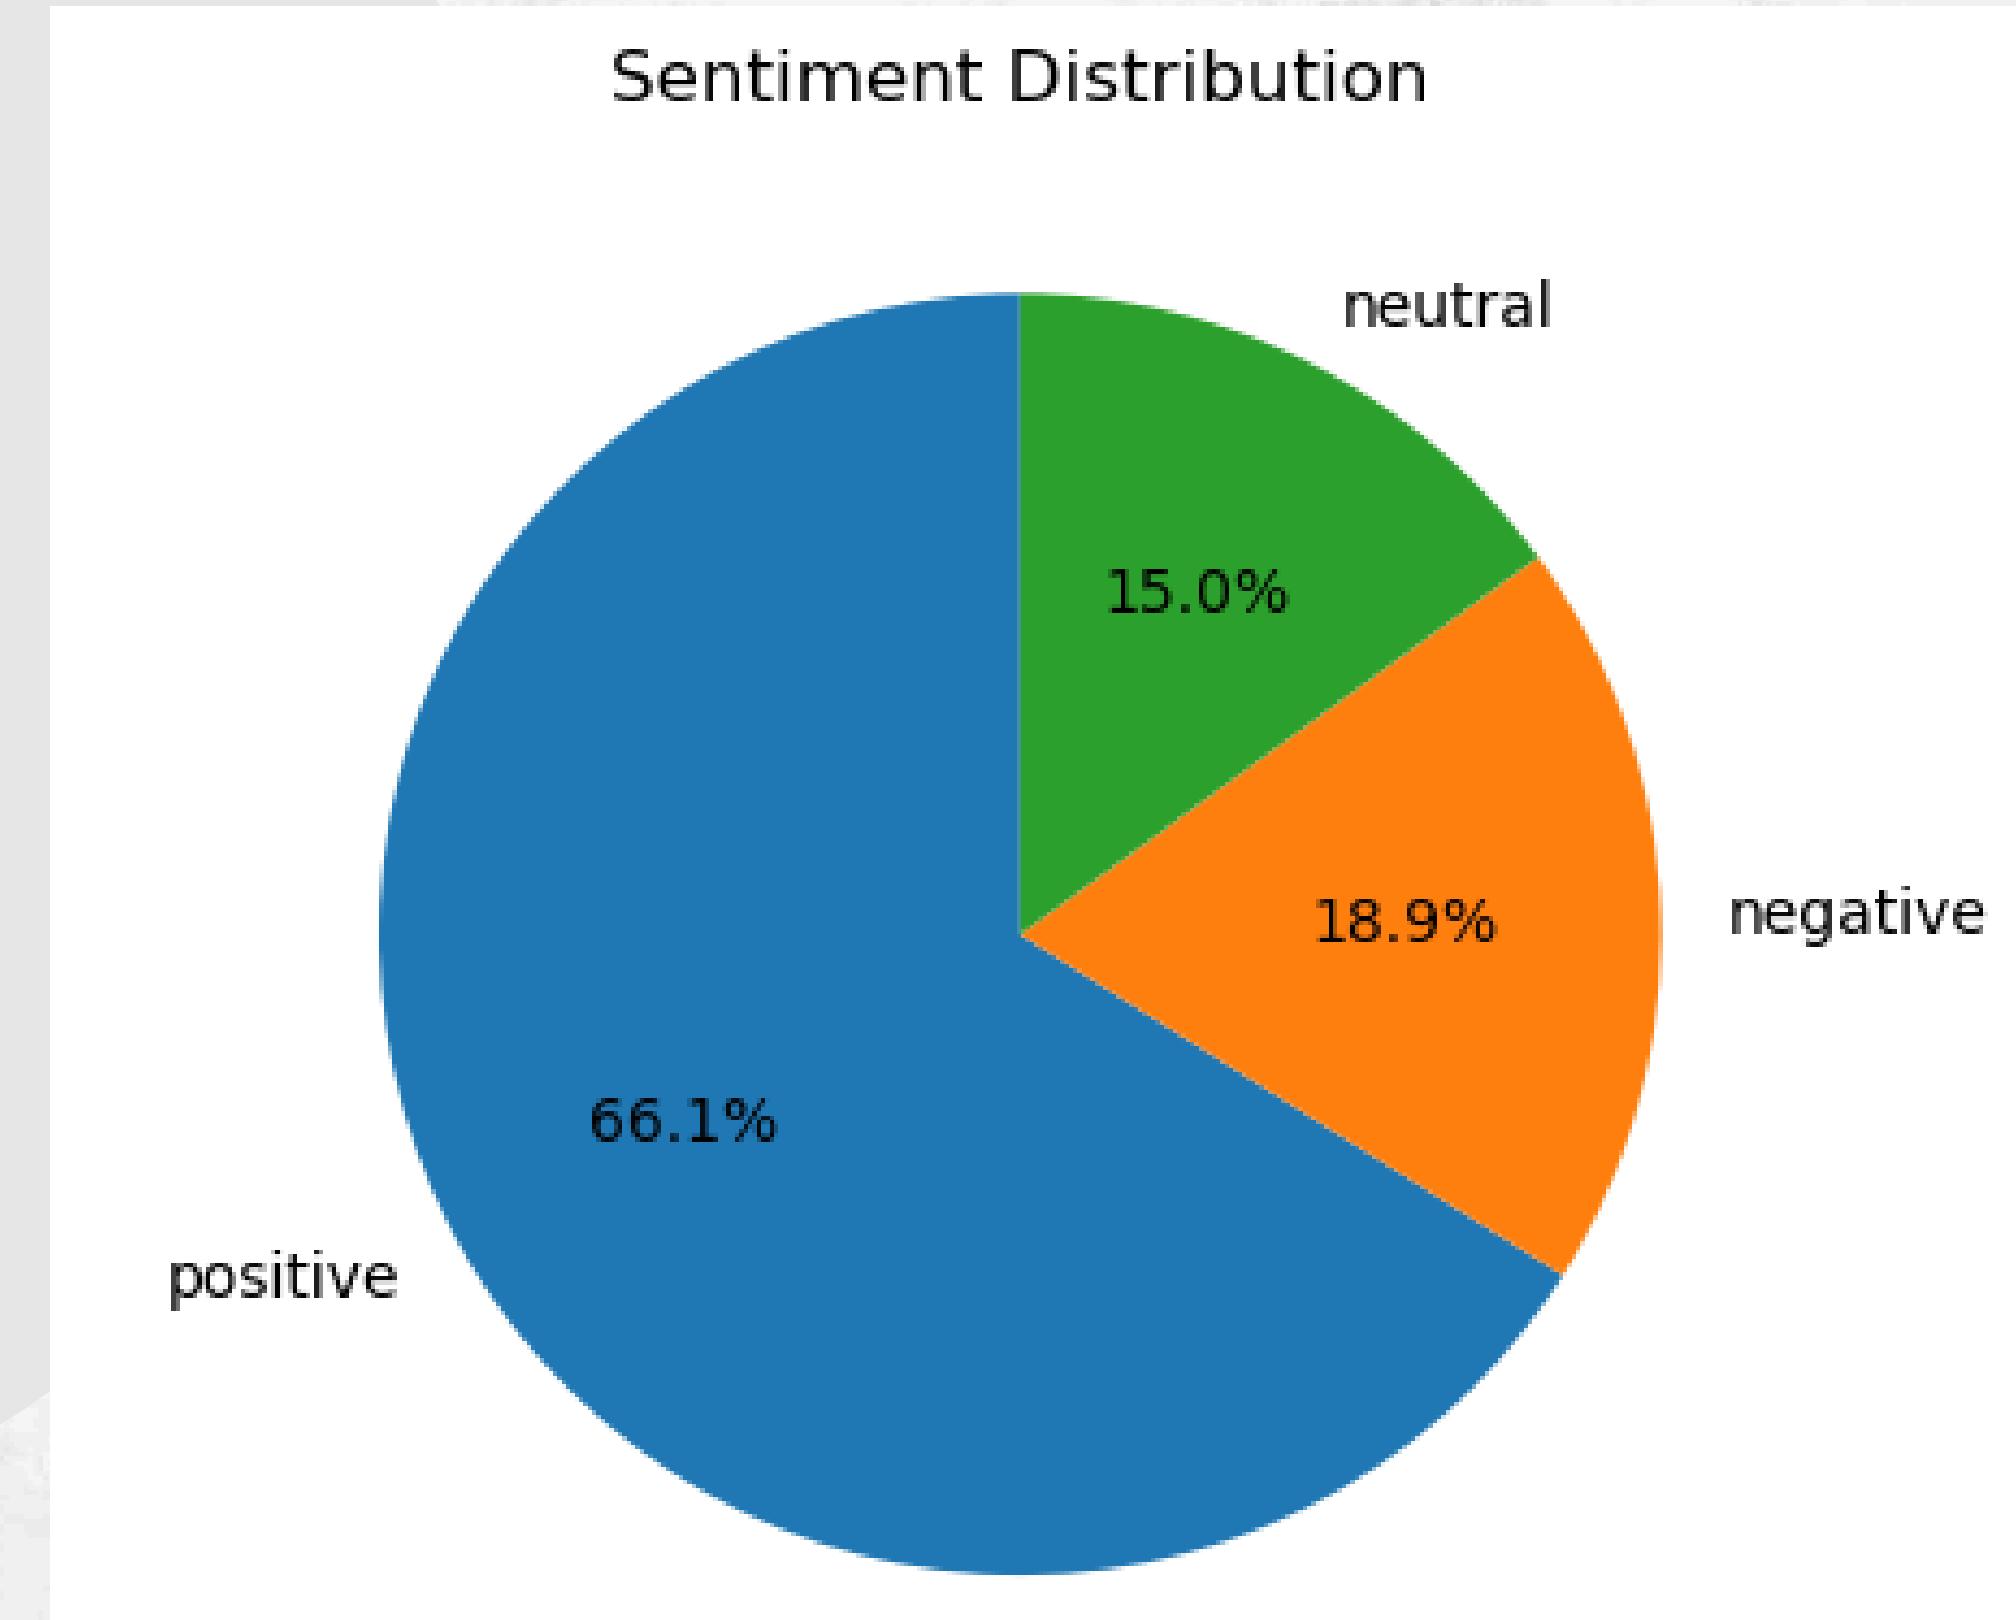
crab legs	22.400406
prime rib	13.396332
king crab	12.204996
wicked spoon	10.685458
food good	9.708715
quality food	9.376912
wait line	9.035394
las vegas	8.776860
buffet vegas	8.705305
bacchanal buffet	8.067720

TRIGRAMS

WORD COUNTS

king crab legs	5.302003
best buffet vegas	3.865627
snow crab legs	3.303073
don waste money	2.589910
buffet las vegas	2.587586
got food poisoning	2.440797
definitely worth price	2.202085
waste time money	1.977260
wicked spoon better	1.966614
don waste time	1.850055

Conclusion



Thanks



0530 308 63 34



afizgordu@gmail.com